


Operating Systems
(INFR09047)
2019/2020 Semester 2

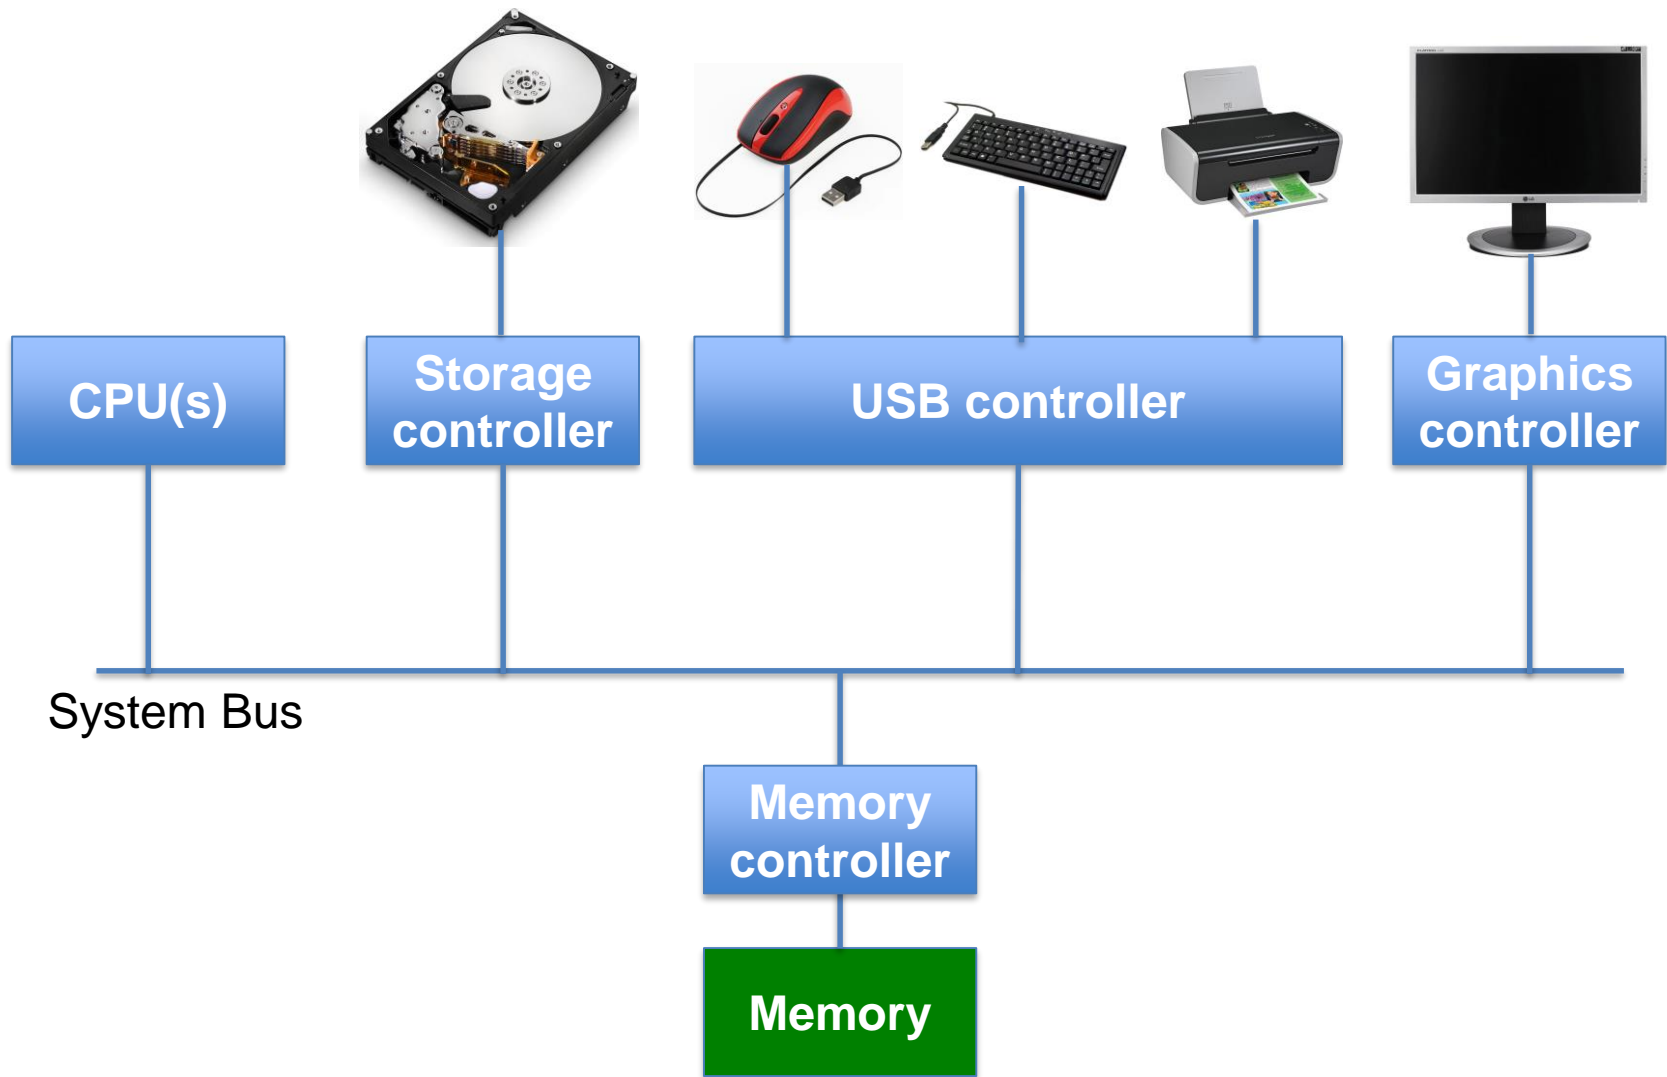
Introduction

abarbala@inf.ed.ac.uk

Computing systems are everywhere



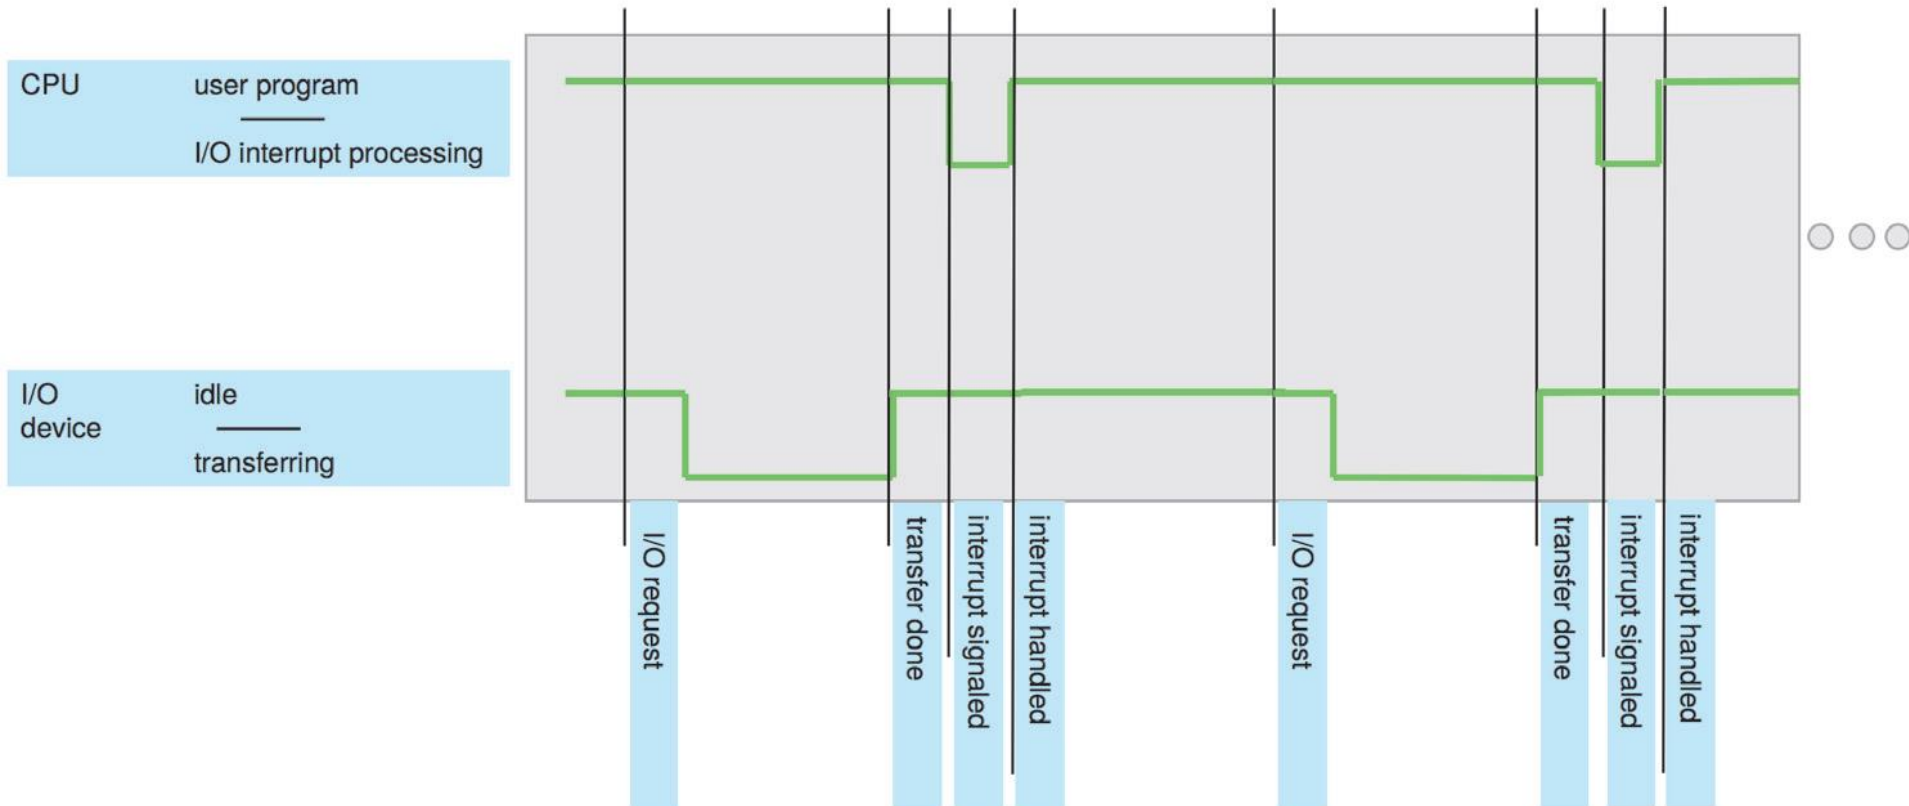
Modern computer system



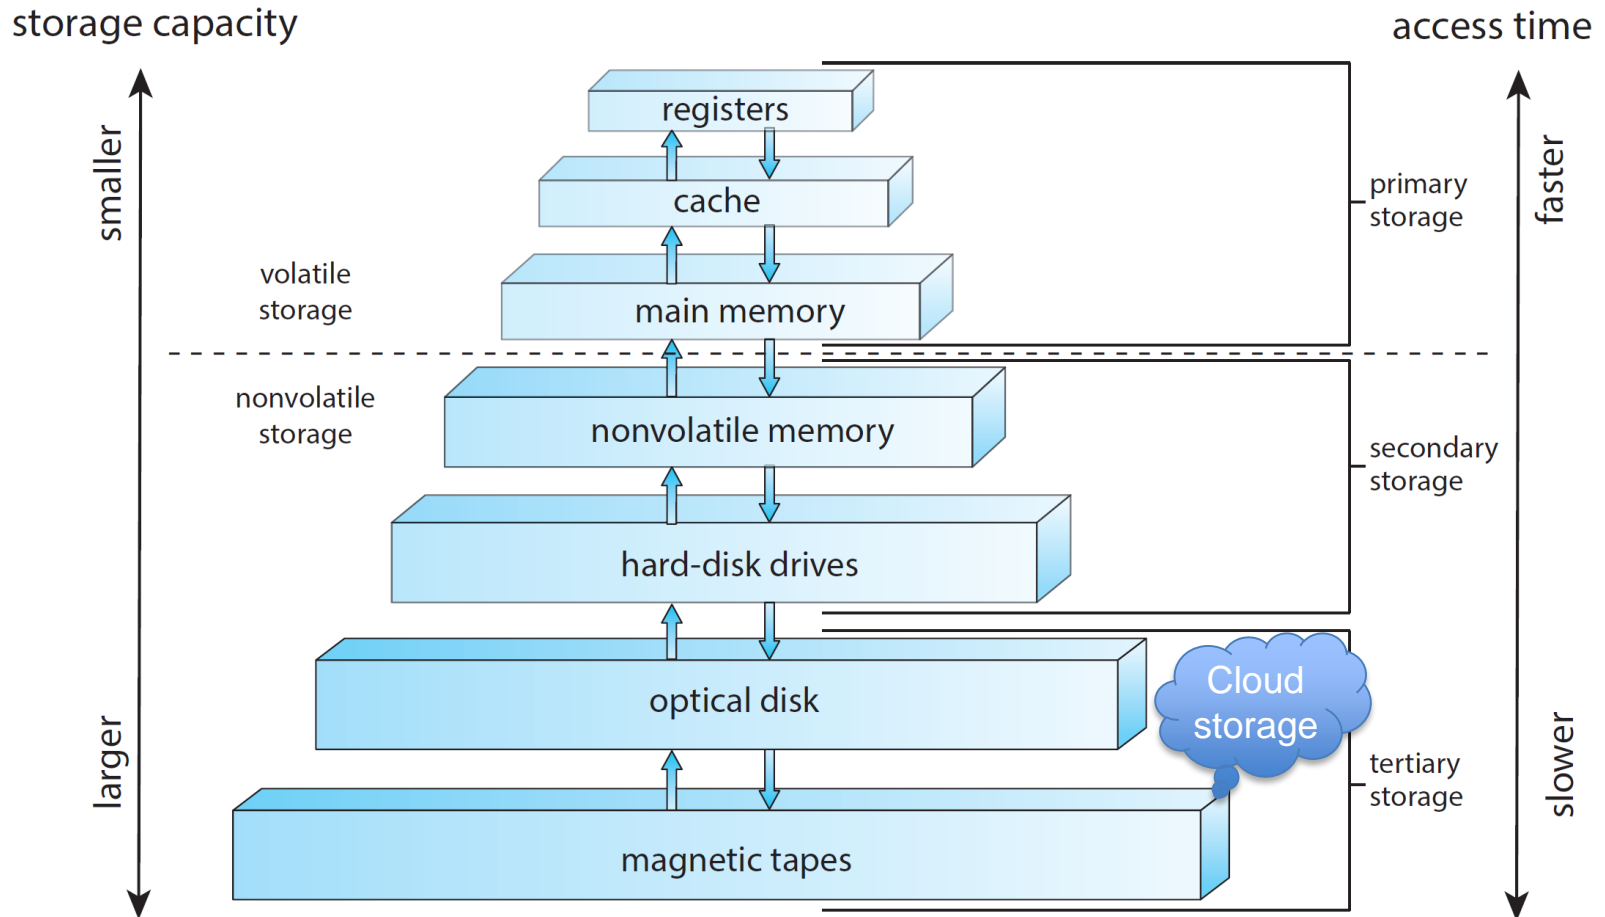
What is an Operating System?

- A program that **manages** a computer's hardware
- A program that acts as an **intermediary** between the user of a computer and computer hardware
- A big program
 - “The **Linux Kernel** Enters 2020 At **27.8 Million Lines** In Git But With Less Developers For 2019”, 1 January 2020 at 09:14 AM EST
 - https://www.phoronix.com/scan.php?page=news_item&px=Linux-Git-Stats-EOY2019



Hardware Recap: Interrupts





Hardware Recap: Storage Structure



Hardware Recap: Memory and CPU

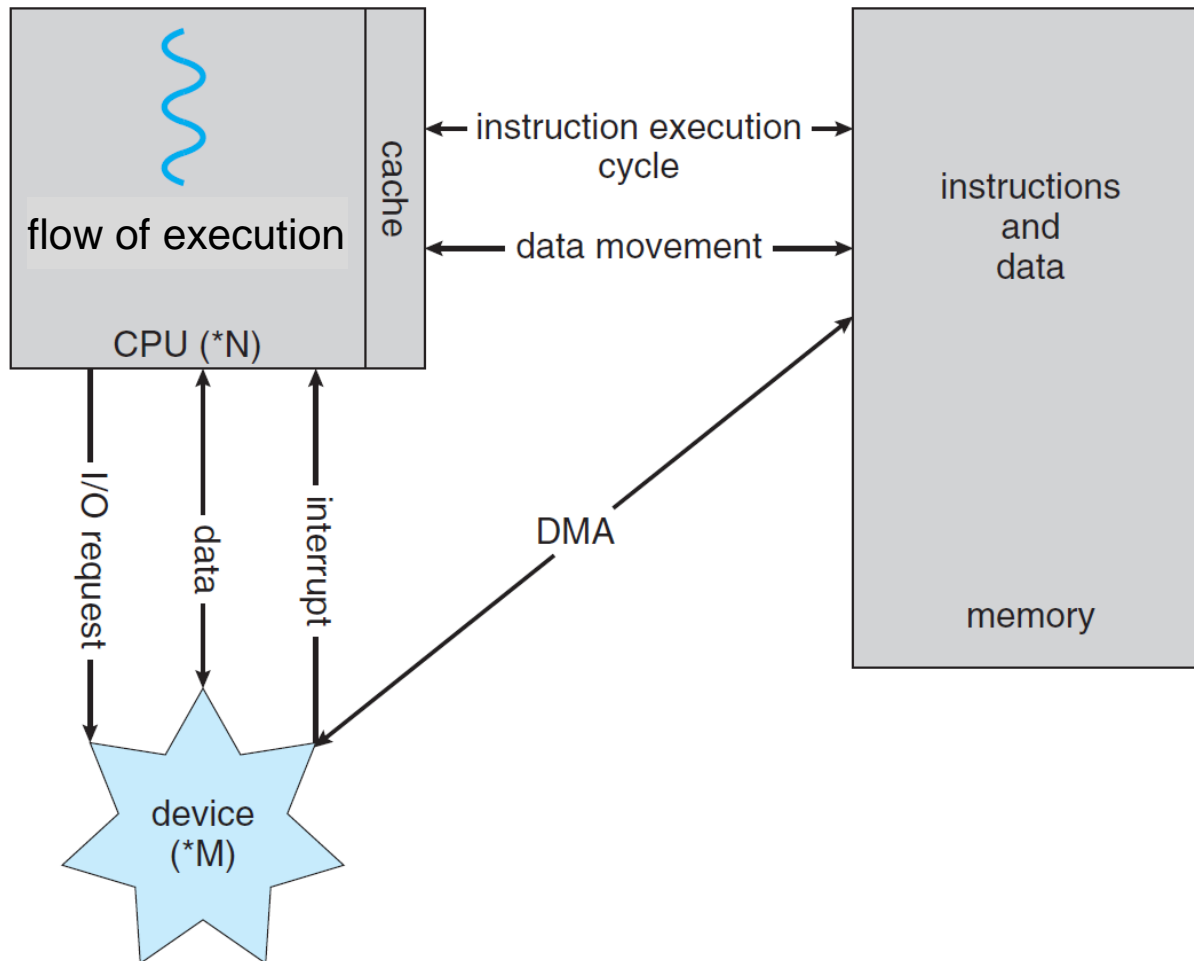


Multicore



NUMA

Hardware Recap: Big Picture



Operating Systems

MS-DOS




MINIX 3

iOS



Mac OS X



Linux™




Some goals of Operating Systems

- Simplify the execution of user programs and make **solving user problems easier**
- Use computer hardware **efficiently**
 - Allow sharing of hardware and software resources
- Make application software **portable and versatile**
- Provide **isolation, security and protection** among user programs
- Improve overall system **reliability**
 - error confinement, fault tolerance, reconfiguration


The traditional Picture

- “The OS is everything you don’t need to write in order to run your application”
- Think of the OS as a library
 - In some ways, it is
 - all operations on I/O devices require OS calls (*syscalls*)
 - In other ways, it isn't
 - you use the CPU/memory without OS calls
 - it intervenes without having been explicitly called

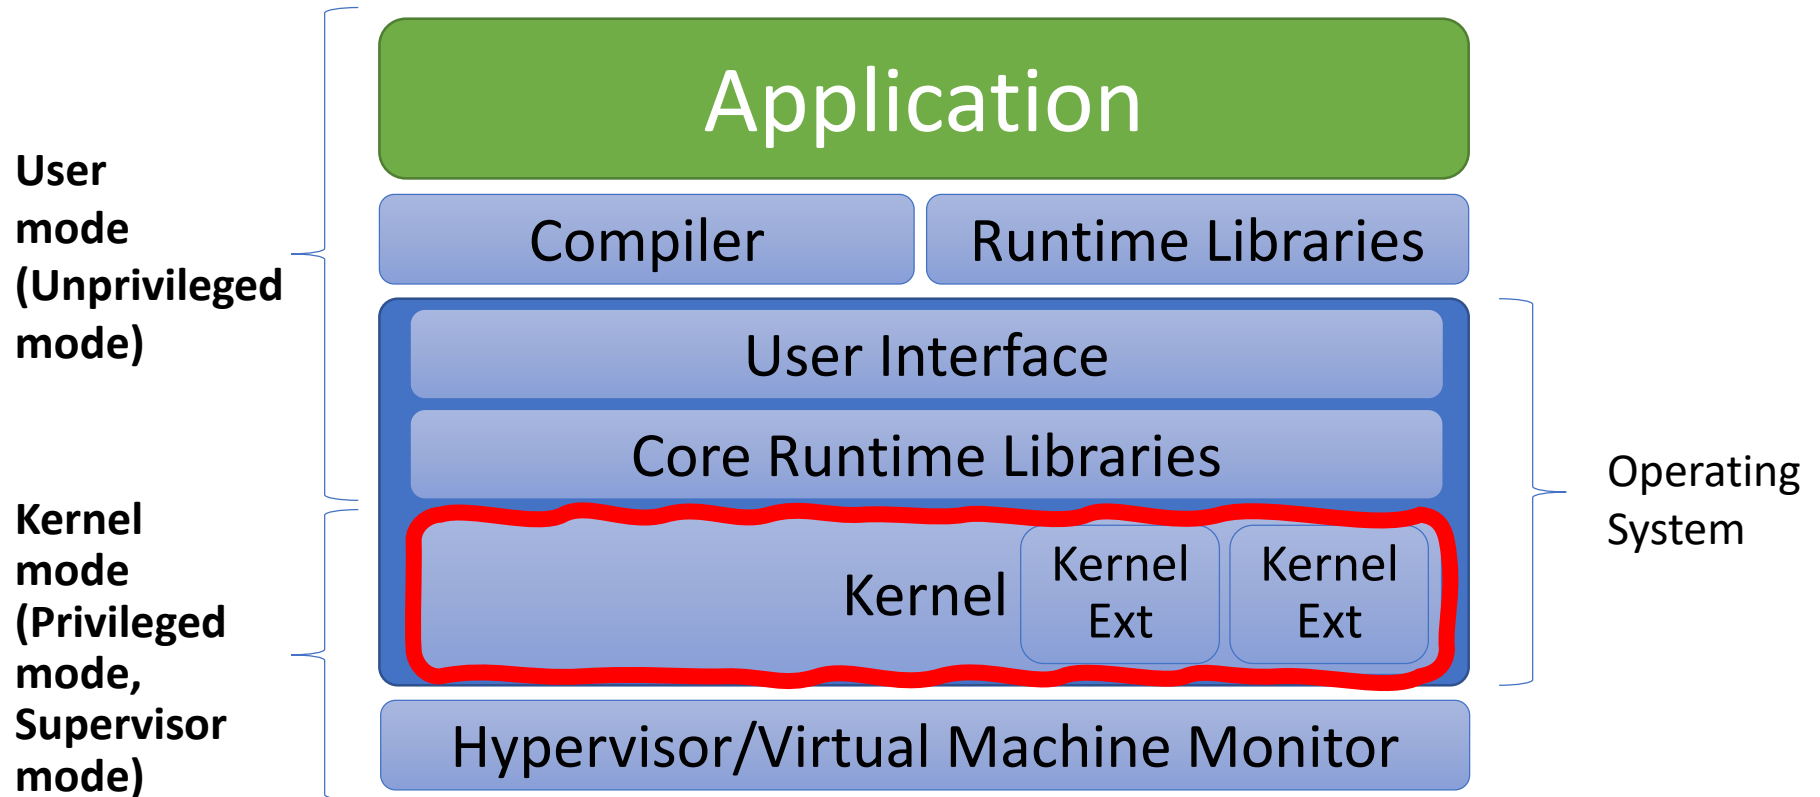


https://en.wikipedia.org/wiki/File:Operating_system_placement.svg

What is OS? #1



What is OS? #2



NOTE there exist OSes that do not use modes, there is hardware that doesn't support modes

The OS and Hardware

- An OS **mediates** programs' access to hardware resources (*sharing and protection*)
 - computation (CPU)
 - volatile storage (memory) and persistent storage (disk, etc.)
 - network communications (TCP/IP stacks, Ethernet cards, etc.)
 - input/output devices (keyboard, display, sound card, etc.)
- The OS **abstracts** hardware into **logical resources** and well-defined **interfaces** to those resources (*ease of use*)
 - processes (CPU, memory)
 - files (disk)
 - sockets (network)

Why Bother with an OS?

- Application benefits
 - programming **simplicity**
 - see high-level abstractions (files) instead of low-level hardware details (device registers)
 - abstractions are **reusable** across many programs
 - **portability** (across machine configurations or architectures)
 - device independence: 3com card or Intel card?
- User benefits
 - **safety**
 - program “sees” its own virtual machine, thinks it “owns” the computer
 - OS **protects** programs from each other
 - OS **fairly multiplexes** resources across programs
 - **efficiency** (cost and speed)
 - **share** one computer across many users
 - **concurrent** execution of multiple programs



Checkpoint

- What is an Operating System?
- What are the benefits of abstraction?
- What other benefits does an OS provide?

Hardware Complexity Increases

Microprocessor Transistor Counts 1971-2011 & Moore's Law


Moore's Law: 2X
transistors/Chip Every 1.5
years (or approx. 2 years)



Transistor: the
basic switching (0/1)
component (discrete)

https://en.wikipedia.org/wiki/Transistor_count
https://en.wikipedia.org/wiki/Moore%27s_law

Software Complexity Increases



Wolverhampton Instrument for Teaching Computing from Harwell (WITCH)



1951
No OS!

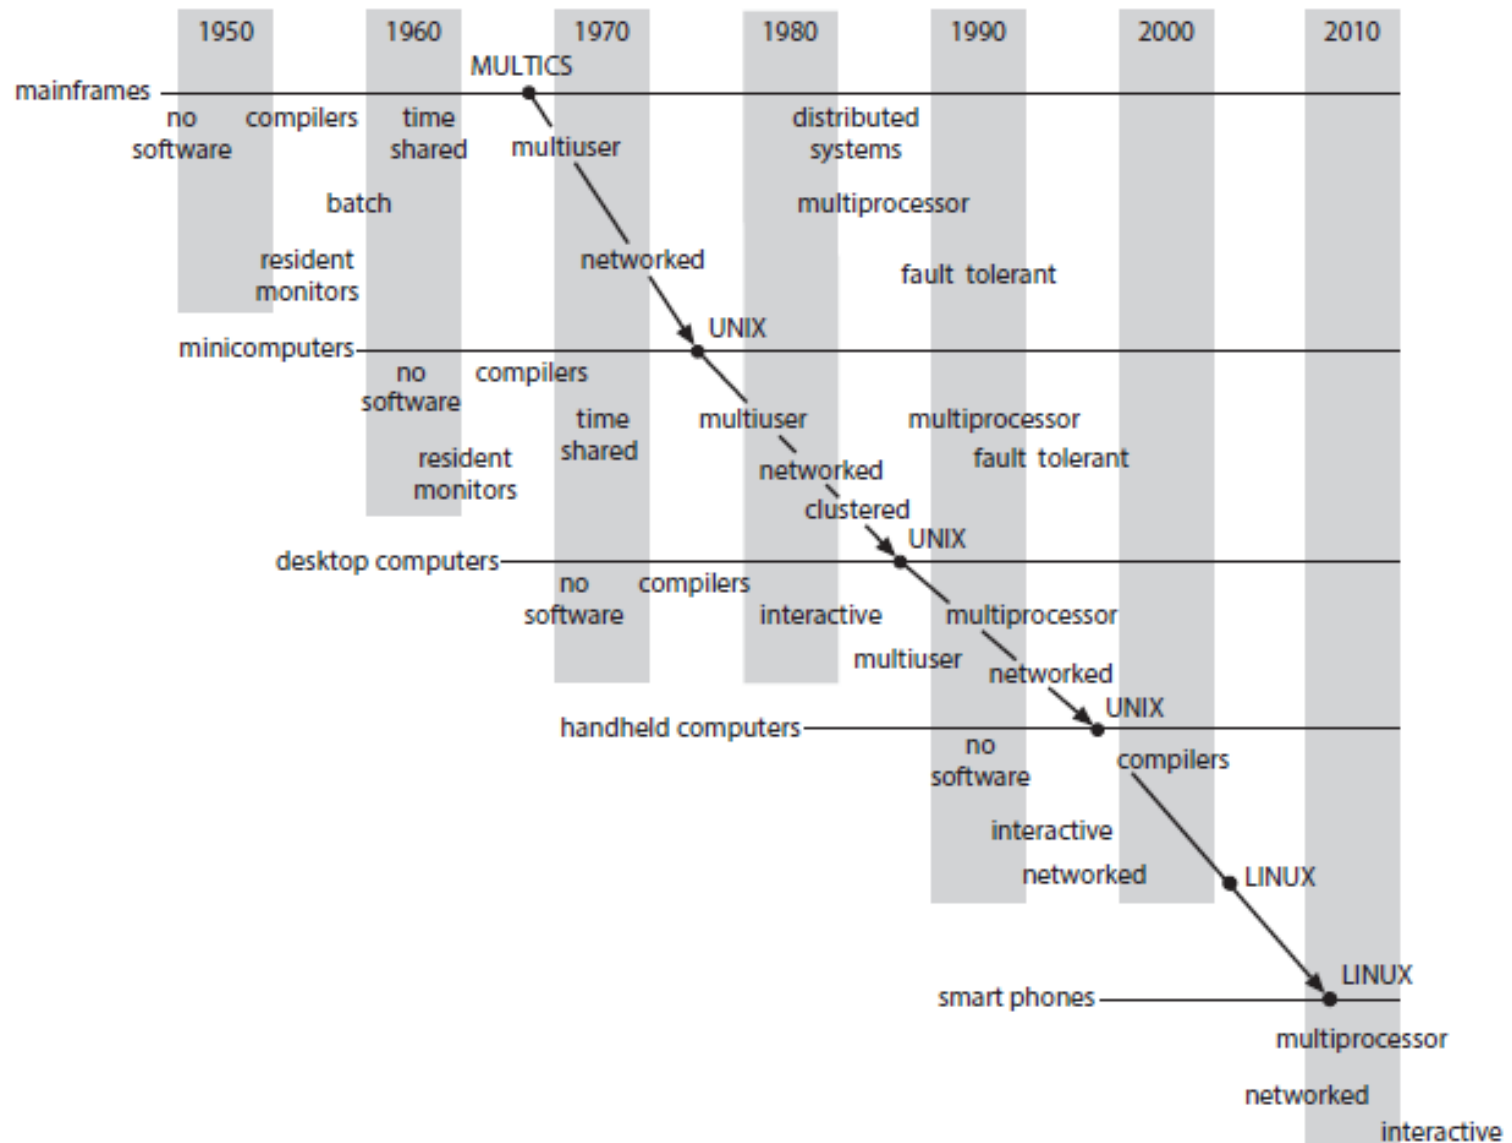


Hardware/Software Changes with Time

- 1960s: mainframe computers (IBM)
- 1970s: minicomputers (DEC)
- 1980s: microprocessors and workstations (SUN), local-area networking, the Internet
- 1990s: PCs (rise of Microsoft, Intel, Dell), the Web
- 2000s:
 - Internet Services / Clusters (Amazon)
 - General Cloud Computing (Google, Amazon, Microsoft)
 - Mobile/ubiquitous/embedded computing (iPod, iPhone, iPad, Android)
- 2010s: sensor networks, “data-intensive computing,” computers and the physical world
- 2020: it’s up to you!!



Progression of Concepts and Form Factors

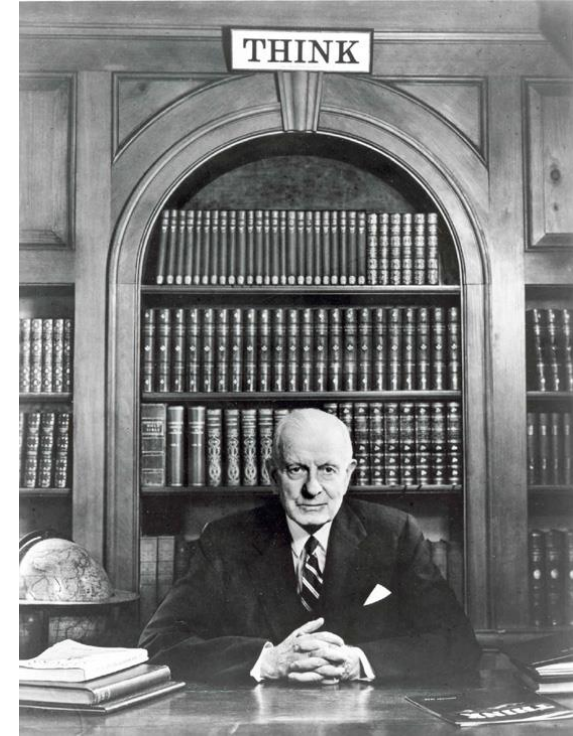


An OS History Lesson


- Operating systems are the result of a 60 year long evolutionary process
 - They were born out of need
- We'll follow a bit of their evolution
- That should help make clear what some of their functions are, and why

In the Beginning...

- 1943
 - T.J. Watson (created IBM):
“I think there is a world market for maybe five computers.”
- Fast forward ... 1950
 - There are maybe 20 computers in the world
 - They were unbelievably expensive
 - Imagine this: machine time is more valuable than person time!
 - Ergo: efficient use of the hardware is paramount
 - Operating systems are born
 - They carry with them the vestiges of these ancient forces



The Primordial Computer




The OS as a linked library

- In the very beginning...
 - OS was just a library of code that you linked into your program; programs were loaded in their entirety into memory, and executed
 - “OS” had an “API” that let you control
 - Disk
 - Printer
 - ...
 - Interfaces were literally switches and blinking lights
 - When you were done running your program, you’d leave and turn the computer over to the next person



Asynchronous I/O


- The disk(osaurus) was really slow
- Add hardware so that the disk could operate without tying up the CPU
 - Disk controller
- Hotshot programmers could now write code that
 - Starts an I/O
 - Goes off and does some computing
 - Checks if the I/O is done at some later time
- Upside
 - Helps increase (expensive) CPU utilization
- Downsides
 - It's hard to get right
 - The benefits are job specific



IBM 1401

Multiprogramming

- To further increase system utilization, **multiprogramming** OSs were invented
 - keeps multiple runnable jobs loaded in memory at once
 - overlaps I/O of one job with computing of another
 - while one job waits for I/O completion, another job uses the CPU
 - Can get rid of asynchronous I/O within individual jobs
 - Life of application programmer becomes simpler; only the OS programmer needs to deal with asynchronous events
 - How do we tell when devices are done?
 - Interrupts
 - Polling
 - What new requirements does this impose?



IBM System 360

Timesharing

- To support interactive use, create a **timesharing OS**
 - multiple terminals into one machine
 - each user has illusion of entire machine to him/herself
 - optimize response time, perhaps at the cost of throughput
- Timeslicing
 - divide CPU equally among the users
 - if job is truly interactive (e.g., editor), then can jump between programs and users faster than users can generate load
 - permits users to interactively view, edit, debug running programs

From MIT ...

- MIT CTSS system (operational 1961) was among the first timesharing systems
 - only one user memory-resident at a time (32KB memory!)
- MIT Multics system (operational 1968) was the first large timeshared system
 - nearly all OS concepts can be traced back to Multics!
 - “second system syndrome”

Multiple Terminals

- In early 1980s, a single timeshared VAX-11/780 ran computing for all of CSE
- A typical VAX-11/780 was 1 MIPS (1 MHz) and had 1MB of RAM and 100MB of disk
- An Apple iPhone 4 is 1GHz (x1000), has 512MB of RAM (x512), and 32GB of flash (x320)



Parallel Systems

- Some applications can be written as multiple parallel **threads** or **processes**
 - can speed up the execution by running multiple threads/processes simultaneously on multiple CPUs [Burroughs D825, 1962]
 - need OS and language primitives for dividing program into multiple parallel activities
 - need OS primitives for fast communication among activities
 - degree of speedup dictated by communication/computation ratio
 - many flavors of parallel computers today
 - SMPs (symmetric multi-processors)
 - MPPs (massively parallel processors)
 - NOWs (networks of workstations)
 - Massive clusters (Google, Amazon.com, Microsoft)
 - Computational grid (SETI @home)

Personal Computing

- Primary goal was to enable new kinds of applications
- Bit mapped display [Xerox Alto, 1973]
 - new classes of applications
 - new input device (the mouse)
- Move computing near the display
 - why?
- Window systems
 - the display as a managed resource
- Local area networks (Ethernet)
 - why?
- Effect on OS?



Distributed OS

- Distributed systems to facilitate use of geographically distributed resources
 - workstations on a LAN
 - servers across the Internet
- Supports applications running among multiple computers
 - interprocess communication
 - message passing, shared memory
 - networking stacks
- Sharing of distributed resources (hardware, software)
 - load balancing, authentication and access control, ...

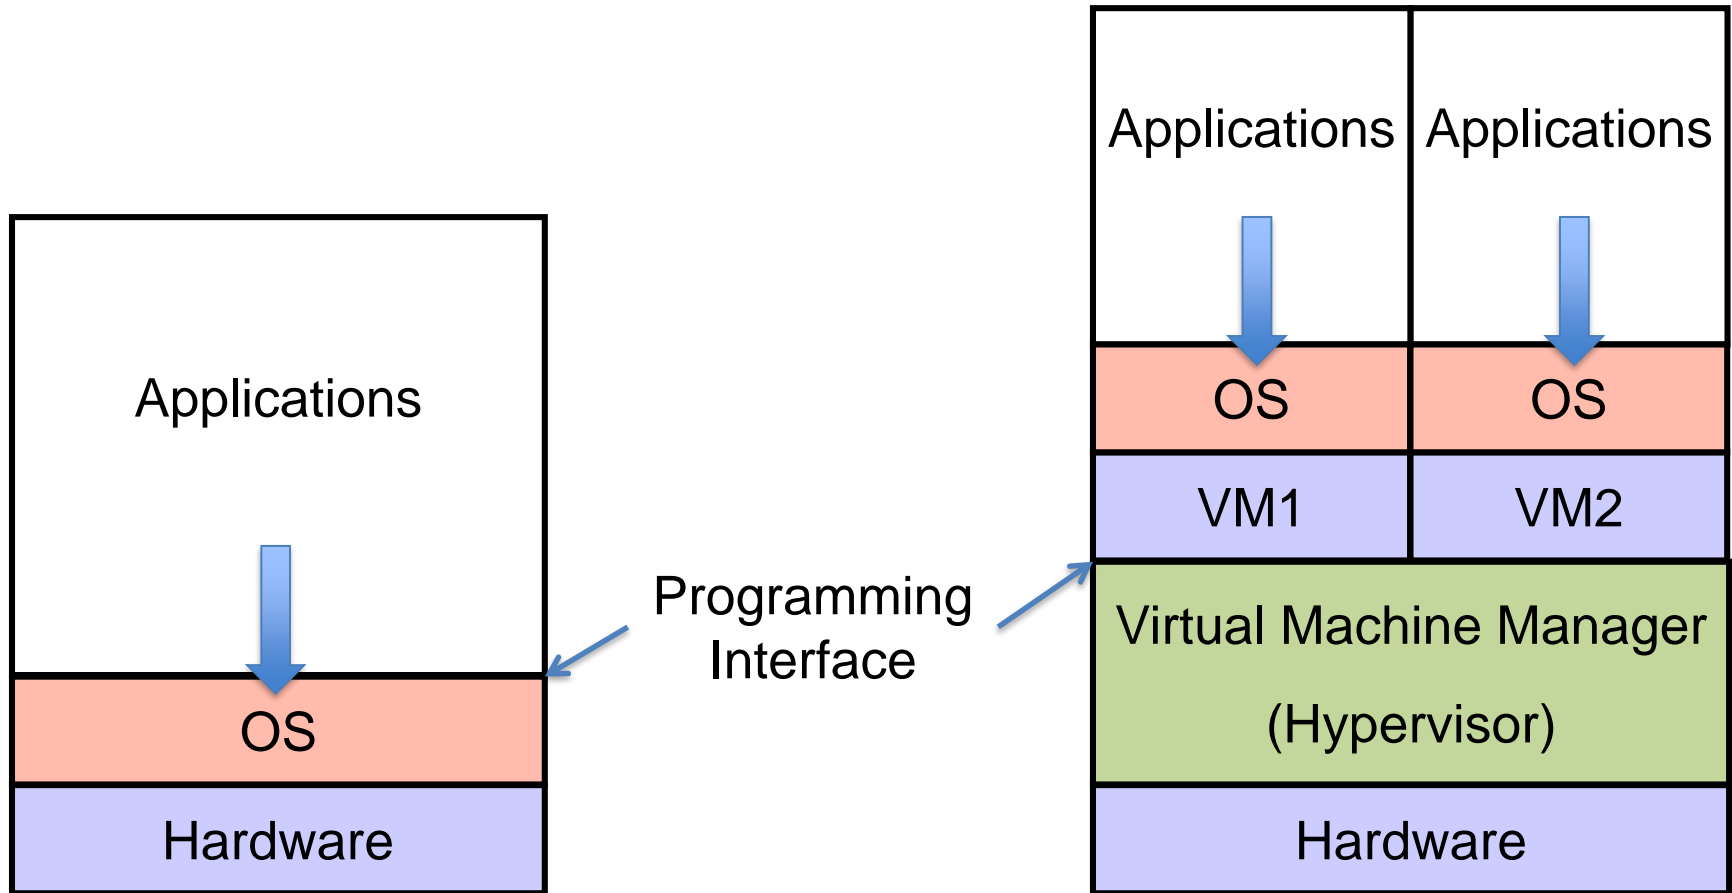
Client/Server Computing

- Mail server/service
- File server/service
- Print server/service
- Compute server/service
- Game server/service
- Music server/service
- Web server/service
- etc.

Peer-to-Peer (p2p) Systems


- Napster
- Gnutella
- BitTorrent
 - example technical challenge: self-organizing overlay network
 - technical advantage of BitTorrent?
 - er ... legal advantage of BitTorrent?

Virtualization



Cloud Computing

- Amazon EC2
- Microsoft Azure
- HP Helion Public Cloud
- Huawei Cloud
- ...



The major OS issues

- **structure**: how is the OS organized?
- **sharing**: how are resources shared across users?
- **naming**: how are resources named (by users or programs)?
- **security**: how is the integrity of the OS and its resources ensured?
- **protection**: how is one user/program protected from another?
- **performance**: how do we make it all go fast?
- **reliability**: what happens if something goes wrong (either with hardware or with a program)?
- **extensibility**: can we add new features?
- **communication**: how do programs exchange information, including across a network?

More OS issues...

- **concurrency**: how are parallel activities (computation and I/O) created and controlled?
- **scale**: what happens as demands or resources increase?
- **persistence**: how do you make data last longer than program executions?
- **distribution**: how do multiple computers interact with each other?
- **accounting**: how do we keep track of resource usage, and perhaps charge for it?

There are tradeoffs, not right and wrong!

Why Should One Learn Operating Systems?

- You may not ever build an OS
- Almost all code runs on top an OS
 - Almost every computing device runs an OS
- But you need to understand the foundations
 - As a Computer Scientist or Computer Engineer
- Knowledge of how OSes work is crucial to proper, efficient, effective, and secure programming

It is a Great Time to Study Operating Systems!

- Many open-source OS projects
 - Check, study, and modify the code
 - Help find and fix bugs
- Emulators to run and debug OSes
 - Existent OSes
 - Historical/Future OSes (hardware not available)
- Renewed interest in OSes from major IT Companies
 - Rise of unikernels
 - Fully formally verified OSes
 - New OSes (e.g. Google Fuchsia)