


Operating Systems
(INFR09047)
2019/2020 Semester 2

Scheduling


abarbala@inf.ed.ac.uk

How Many Processes or Threads Are Running On My Computer?



Window's Task Manager

UNIX's htop (or top)

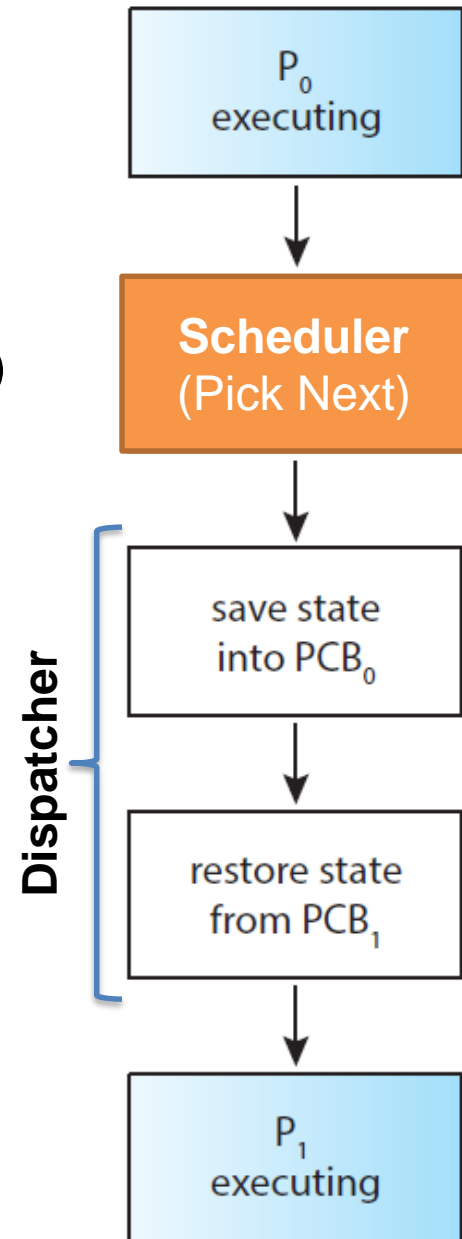


Who will run next?

Scheduling

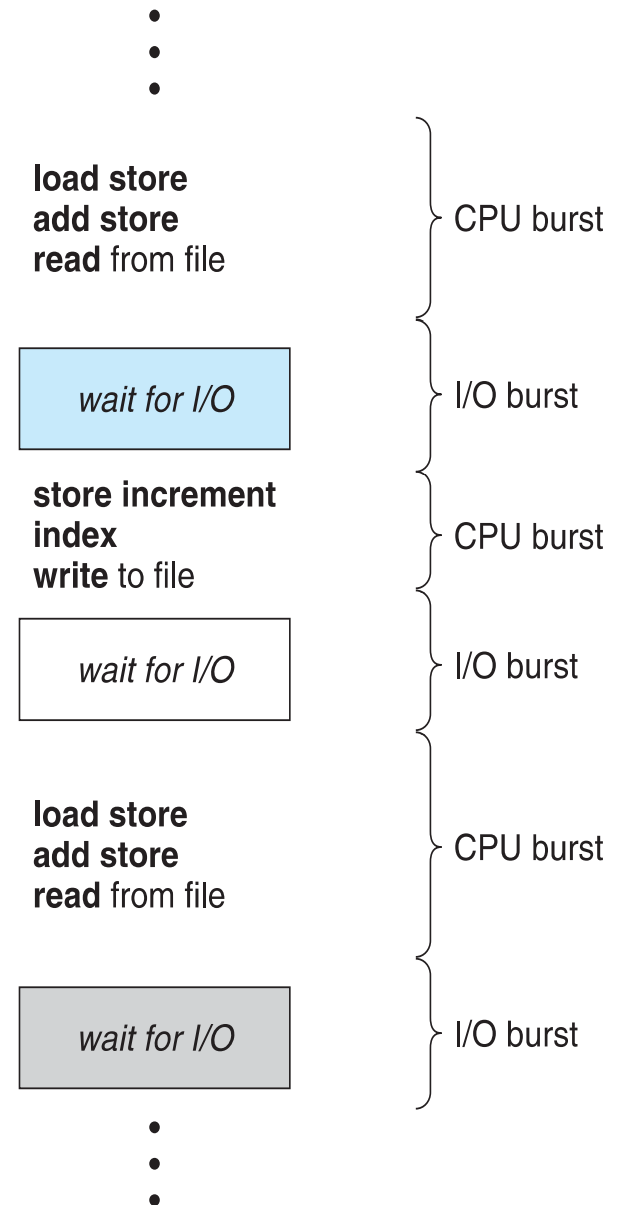
Single CPU/core

- Decision of who will **run next**
 - Process or threads
 - **In kernel for kernel-level threading** (slideset' focus)
- **Pick among** the ones in ready queue(s)
- *When?*
 - Potentially, **all times we switch** to the OS
 - Interrupt (device completion, timer interrupt, etc.)
 - Syscall (including voluntary process/thread termination, or yield)
 - Exception (including involuntary termination)
- *How?*
 - **Scheduler** decides
 - **Policy** (implemented by an algorithm)
 - Task switched by **dispatcher**
 - **Mechanism** (kernel or user code)



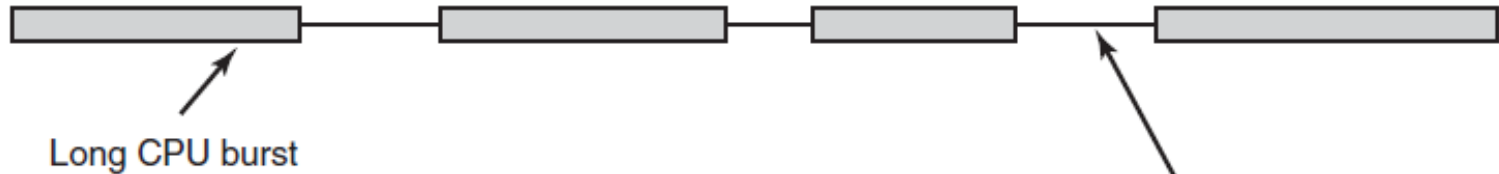
Process or Thread Behavior #1

- Process or thread execution consists of cycles of
 - CPU execution: **CPU burst**
 - I/O wait: **IO burst**
- **CPU bursts** distribution is **application dependent**
- **Maximum CPU utilization** with multiprogramming
 - Don't leave CPU idle
- While a process or thread is waiting for IO **another can run** on the CPU
 - Focus on **single** CPU/core



Process or Thread Behavior #2

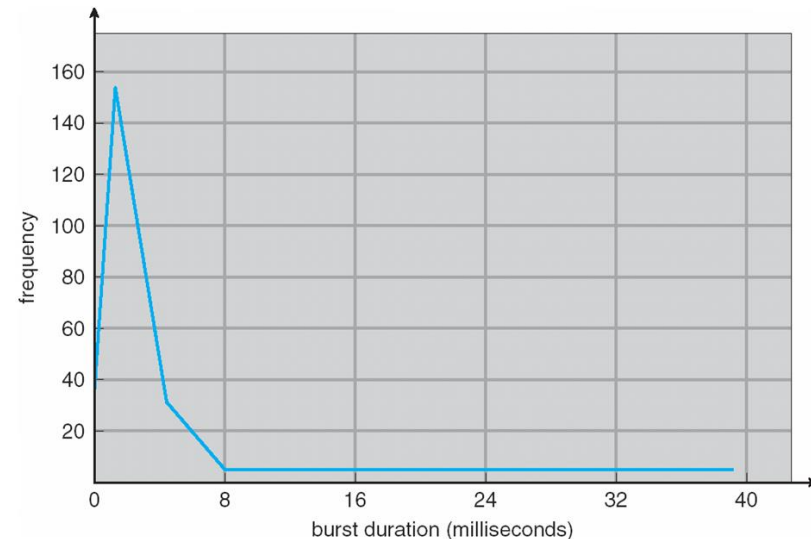
CPU
bound



I/O
bound



Time



While a process or thread is waiting for I/O **another can run** on the CPU

Scheduling Goals: Performance

- Many performance goals (which may conflict)
 - Maximize **CPU utilization**
 - Maximize **throughput** (processes completed per time unit)
 - Minimize **turnaround time** (time from submission of task to completion)
 - Minimize **waiting time** (all periods spent waiting in the ready queue from submission)
 - Minimize **response time** (time from submission of request to response is produced)
 - Minimize **energy** (joules per instruction) subject to some constraint (e.g., frames/second)
- In most cases we optimize the **average metric**
 - But sometimes minimize the worst case (e.g., response time)

Scheduling Goals: Fairness

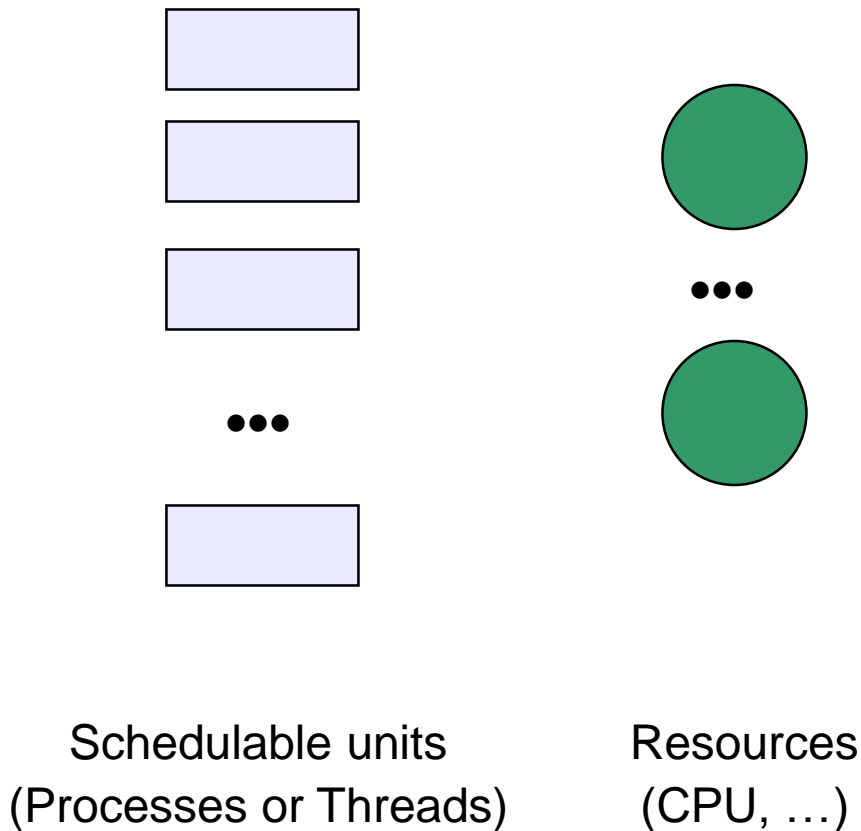
- No single, compelling definition of **“fair”**
 - How to **measure** fairness?
 - Equal CPU consumption? (over what time scale?)
 - Fair per-user? per-process? per-thread?
 - What if one process is CPU bound and one is I/O bound?
- Sometimes the goal is to be **unfair**
 - Explicitly favor some particular class of requests
 - **Priority system**
 - Avoid starvation
 - Be sure everyone gets at least some service

Classes of Schedulers

- Batch
 - Throughput / utilization oriented
 - *Example:* audit inter-bank funds transfers each night, Pixar rendering, Hadoop/MapReduce jobs
- Interactive
 - Response time oriented
 - *Example:* window-based operating system
- Real time
 - Deadline driven
 - *Example:* embedded systems (cars, airplanes, etc.)

We will **not talk** about real-time scheduling

General Scheduling Problem




Scheduling:

- **Who** to *assign* each resource to
- **When** to re-evaluate your decisions

Scheduling is **not** just about assigning *processes and threads* to the CPU, but to **any other HW/SW resource** of the computer

When to Re-evaluate the Decision?

- **Non-preemptive scheduling**
 - Processes/threads execute until completion or until they want
 - **Voluntary process switch**
 - **Process/thread switch on blocking calls**
 - The scheduler gets involved only at exit or on request
 - For every clock interrupt, running process keeps going
- **Preemptive scheduling**
 - While a process/thread executes, its execution may be paused, and another process/thread resumes its execution, etc.
 - **Involuntary process switch**
 - For every clock interrupt, running process may be suspended and switched with another process (if there is any)



wait()s removed to ease the draw

First-come First-served (FCFS)

- Processes are assigned to the CPU in the **order they request it** (or they **arrive**)
- **Non-preemptive**
- Real-world scheduling of people in (single) lines
 - Supermarkets

Algorithm #1

FCFS Example #1

- Arrival **order** for the processes
 - P1, P2, P3
- Turnaround time
 - P1 = 24
 - P2 = 27
 - P3 = 30
- Average turnaround time
 - $(24+27+30)/3 = 27$
- Short process delayed by long process
 - **Convoy effect**

Process	CPUTime
P1	24
P2	3
P3	3

Turnaround Time – time taken by a job to complete after submission

FCFS Example #2

- Arrival **order** for the processes
 - P2, P3, P1
- Turnaround time
 - P1 = 30
 - P2 = 3
 - P3 = 6
- Average turnaround time
 - $(30+3+6)/3 = 13$
- Much better than the previous case

Process	CPUTime
P1	24
P2	3
P3	3

FCFS Drawbacks

- Average response time can be **poor**
 - Short tasks wait behind big ones (convoy effect)
- May lead to poor utilization of **other resources**
 - Poor overlap of CPU and I/O activity
 - Example
 - A CPU-intensive job prevents an I/O-intensive job from a small bit of computation
 - Preventing it from going back and keeping the I/O subsystem busy

Shortest Job First (SJF)

- Associate with each process the **length of its CPU time**
- Use the CPU time length to schedule the process with the **shortest CPU time first**
- Two variations
 - *Non-preemptive* – once CPU is given to the process, it cannot be taken away until completion (or blocking)
 - *Preemptive* – if a new process arrives with CPU time less than the remaining time of current executing process, preempt.
 - ***Shortest Remaining Time Next, SRTN***

Non-Preemptive SJF Example

- Arrival time for the processes
 - P1 at 0, P2 at 2, P3 at 4, P4 at 5
- Turnaround time
 - $P1 = 7$
 - $P2 = 12 - 2 = 10$
 - $P3 = 8 - 4 = 4$
 - $P4 = 16 - 5 = 11$
- Average turnaround time
 - $(7+10+4+11)/4 = 8$
- 3 ctx switches


Process	ArrivalTime	CPUTime
P1	0	7
P2	2	4
P3	4	1
P4	5	4

Algorithm #2

Preemptive SJF Example

- Arrival time for the processes
 - P1 at 0, P2 at 2, P3 at 4, P4 at 5
- Turnaround time
 - $P1 = 16$
 - $P2 = 7 - 2 = 5$
 - $P3 = 5 - 4 = 1$
 - $P4 = 11 - 5 = 6$
- Average turnaround time
 - $(16+5+1+6)/4 = 7$
- 5 ctx switches

Process	ArrivalTime	CPUTime
P1	0	7
P2	2	4
P3	4	1
P4	5	4



SJF Drawbacks

- Preemptive SJF is **optimal**
- But it can only be approximated
 - **Too complex** to be implemented in practice
 - **Not always possible** to determine the CPU/IO burst

Algorithm #3

Round-robin (RR)


- Each process is allowed to run for a specified time interval
 - Called **quantum**
- After this time **has elapsed**
 1. The process is preempted
 2. And added to the end of the ready queue
 3. The next process is scheduled
- If the process **terminates or blocks for IO** before this time
 1. It is added to a wait queue
 2. The next process is scheduled

Algorithm #3

Round-robin Example

Process	CPU time
P1	53
P2	8
P3	68
P4	24

Time Quantum = 20



- Waiting time for
 - $P_1 = (68 - 20) + (112 - 88) = 72$
 - $P_2 = (20 - 0) = 20$
 - $P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$
 - $P_4 = (48 - 0) + (108 - 68) = 88$
- Average waiting time $(72 + 20 + 85 + 88) / 4 = 66.25$
- Average turnaround time $(125 + 28 + 153 + 112) / 4 = 104.5$

What About the Time Quantum?

- **Context switching** may impact the choice of the time quantum
- Example
 - Context switch is 1ms
 - Time quantum is 4ms
 - 20% of the time is thrown away context switching
- Typical numbers
 - Context switch is in the order of tens of us
 - Timeslice/quantum is 1KHz (every 1ms)
- But when there are a lot of processes a long time quantum causes a **poor response time**

What About Round-robin?


- Advantages
 - Solution to fairness and starvation
 - Fair allocation of CPU across jobs
 - Low average waiting time when job lengths vary
 - Good for responsiveness (interactivity) if small number of jobs
- Disadvantages
 - Context-switching time may add up for long jobs

FCFS vs RR

- Assuming zero-cost context-switching time, is RR better than FCFS?
- Simple example
 - 10 jobs, each takes 100s of CPU time
 - RR scheduler quantum of 1s
 - All jobs start at the same time
- Turnaround Times
 - Both RR and FCFS finish at the same time
 - Average turnaround time is much worse under RR
 - Bad when all jobs same length
- Cache state must be shared between all jobs with (may slow down RR execution)
 - Total time for RR longer even for zero-cost switch

Job #	FIFO	RR
1	100	991
2	200	992
...
9	900	999
10	1000	1000

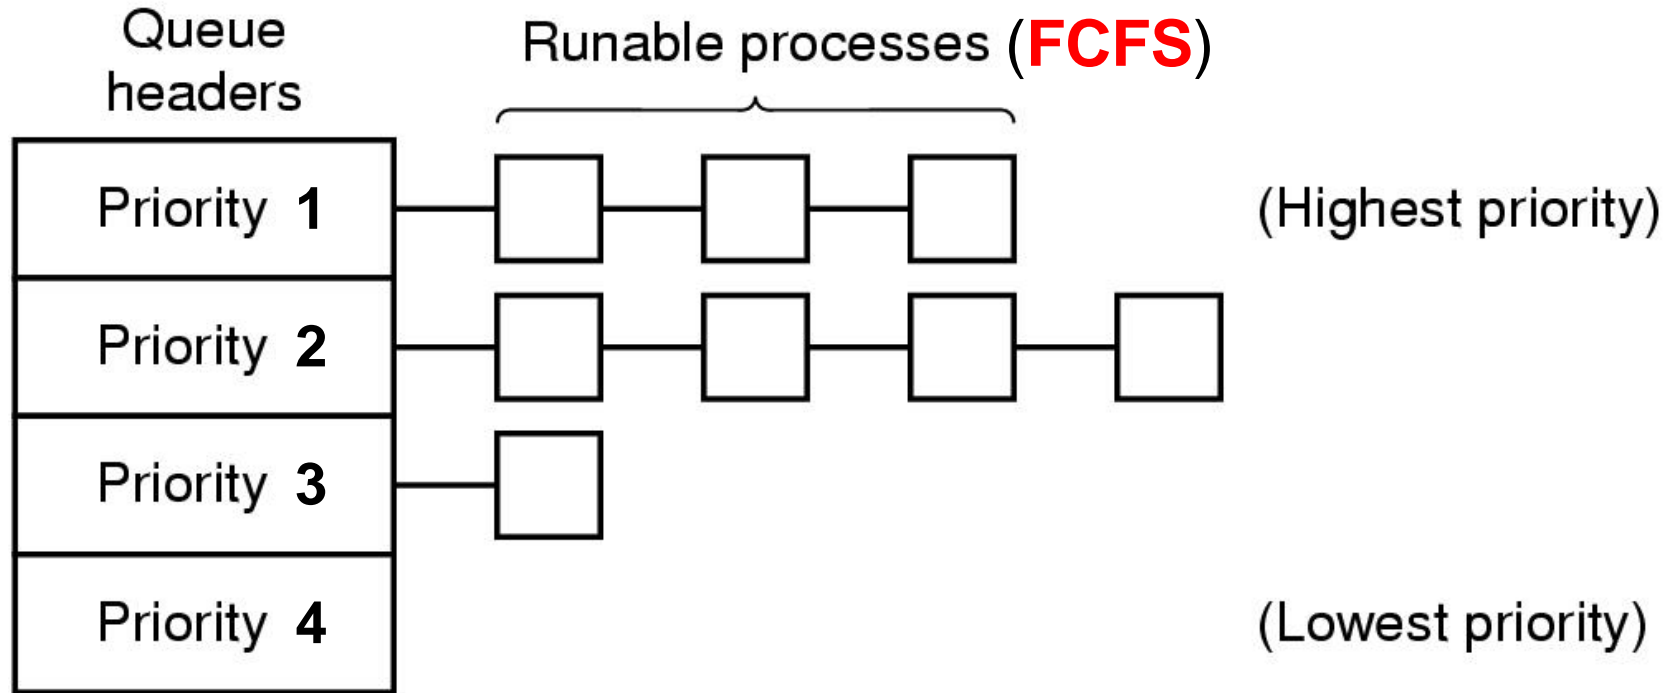
FCFS vs RR with Different Quantum



	Quantum	P ₁	P ₂	P ₃	P ₄	Average
Wait Time	Best FCFS	32	0	85	8	31¼
	Q = 1	84	22	85	57	62
	Q = 5	82	20	85	58	61¼
	Q = 8	80	8	85	56	57¼
	Q = 10	82	10	85	68	61¼
	Q = 20	72	20	85	88	66¼
	Worst FCFS	68	145	0	121	83½
Turnaround Time	Best FCFS	85	8	153	32	69½
	Q = 1	137	30	153	81	100½
	Q = 5	135	28	153	82	99½
	Q = 8	133	16	153	80	95½
	Q = 10	135	18	153	92	99½
	Q = 20	125	28	153	112	104½
	Worst FCFS	121	153	68	145	121¾

Algorithm #4

Priority (PRIO) #1



A scheduling algorithm with four priority classes

Priority (PRIO) #2

- Execution Plan
 - Always execute highest-priority runnable jobs to completion
 - Each queue processed in FCFS
- Problems
 - Starvation
 - Lower priority jobs don't get to run because higher priority tasks always running
 - Deadlock
 - Priority Inversion
 - Not strictly a problem with priority scheduling
 - Happens when low priority task has lock needed by high-priority task (busy waiting)

Algorithm #4

Priority Example

- Turnaround time
 - P1 = 16
 - P2 = 1
 - P3 = 18
 - P4 = 19
 - P5 = 6
- Average turnaround time
 $(16+1+18+19+6)/5 = 12$

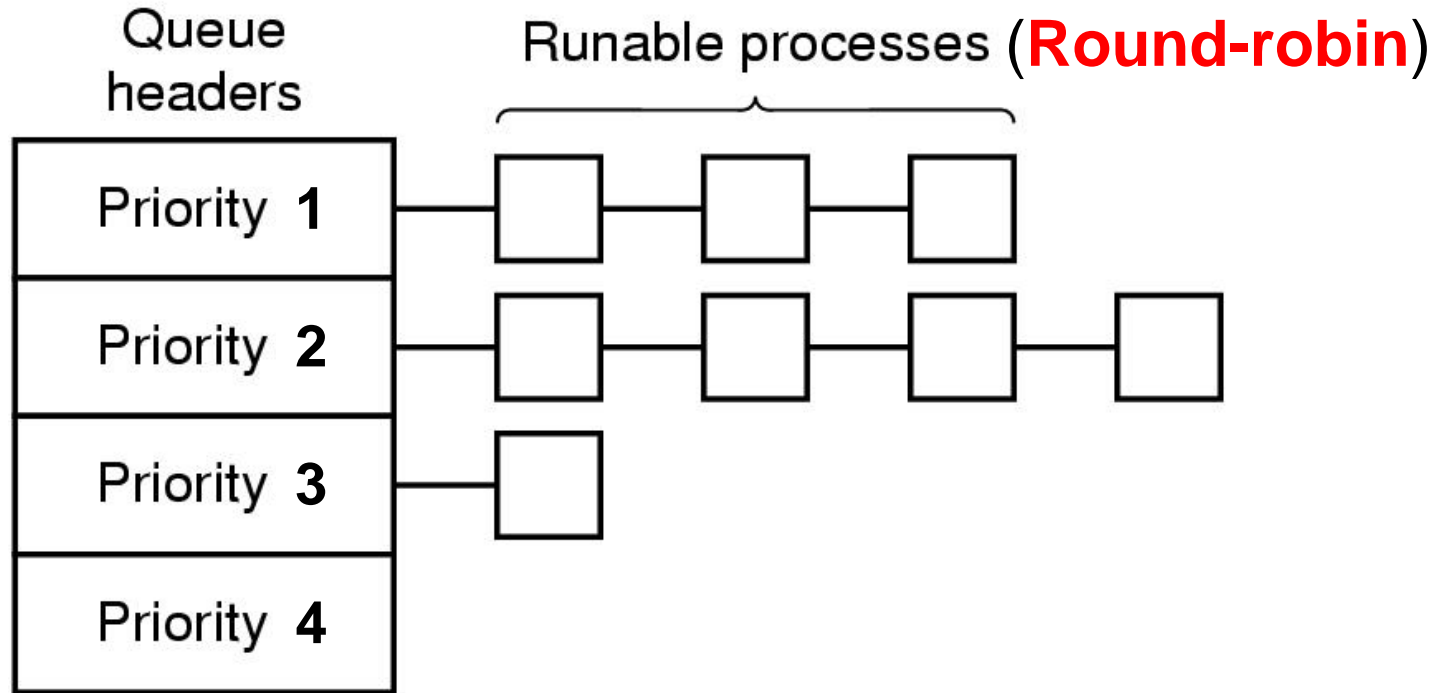
Process	CPUTime	Priority
P1	10	3
P2	1	1
P3	2	4
P4	1	5
P5	5	2

How to Assign Priorities?

- Statically, based on
 - Process type
 - User
 - How much the user paid
- Dynamically, based on how much they run vs IO
 - $\text{Priority} = 1/f$
 - f = size of quantum used last
 - The longer a process ran, the lower its priority
 - The process that runs the shortest gets highest priority to run next

Algorithm #5

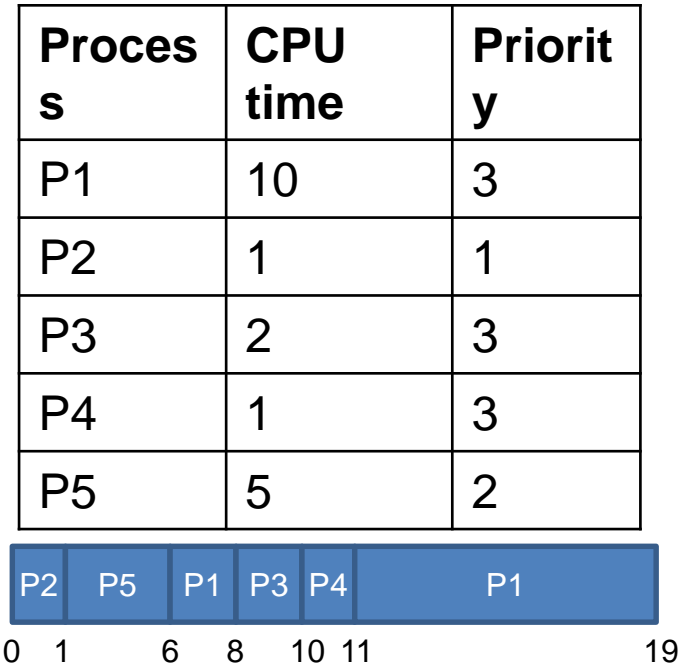
Multiple Queues (MQ)



Algorithm #5

Multiple Queues Example

- Turnaround time
 - P1 = 19
 - P2 = 1
 - P3 = 10
 - P4 = 11
 - P5 = 6
- Average turnaround time
 $(19+1+10+11+6)/5 = 9.4$

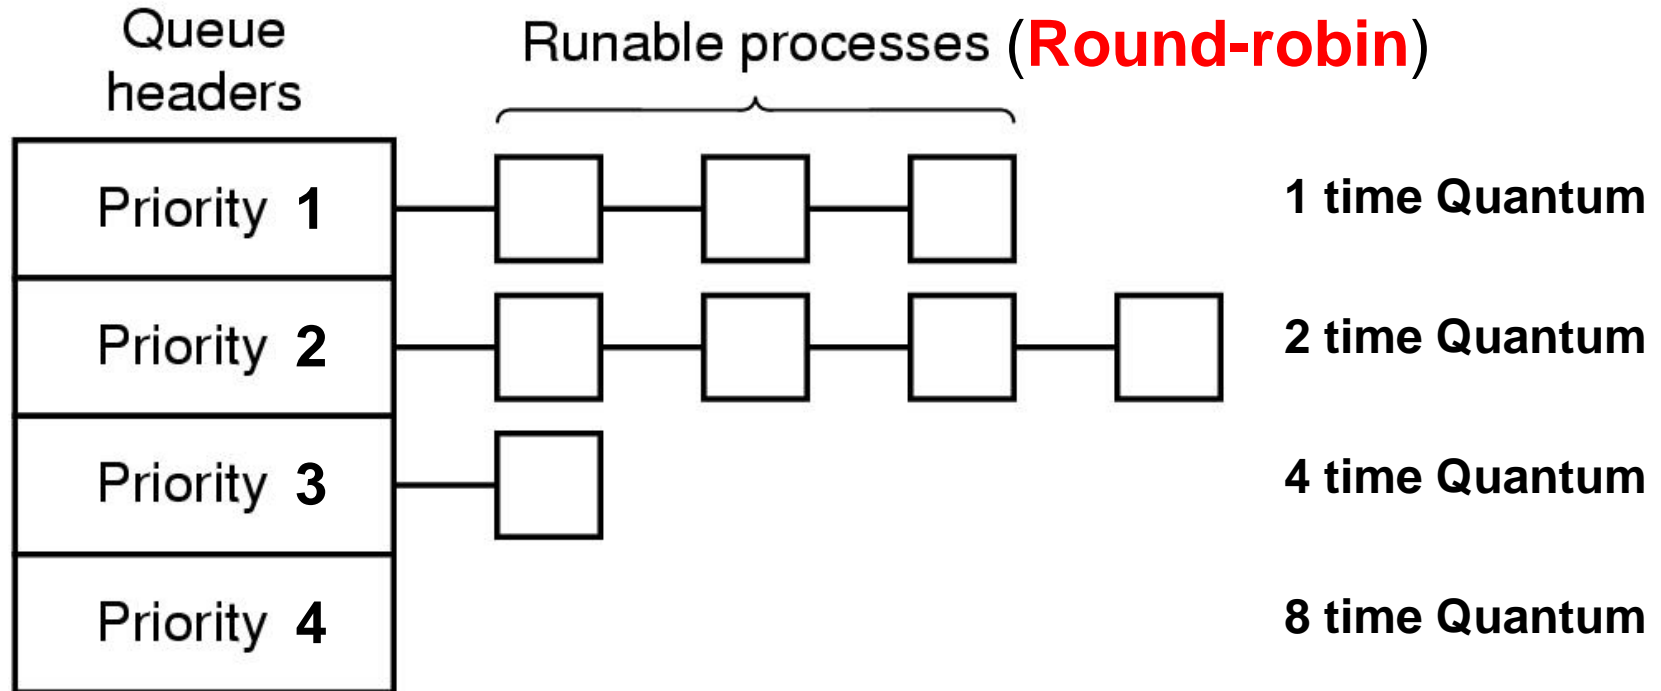


Multilevel Feedback Queue (MLFQ) #1

- Execution Plan
 - Same as MQ scheduling
 - But each queue has a different time quanta
 - Shortest for high-prio
 - Longer for low-prio
 - Processes start at the highest priority
 - When a process **exceeds** its quanta is moved to the lower priority
 - When a process **becomes interactive** is moved to higher priority
- Problem
 - If the user discovers how make his/her tasks interactive he/she can play the system

Algorithm #6

Multilevel Feedback Queue (MLFQ) #2




Algorithm #6

MLFQ Example

- Turnaround time
 - P1 = 102
 - P2 = 3
- Average turnaround time
 $(102+3)/2 = 52.5$
- 8 Context Switches
- vs 101 context switches with **fixed quantum**

Processes	CPU time
P1	100
P2	2



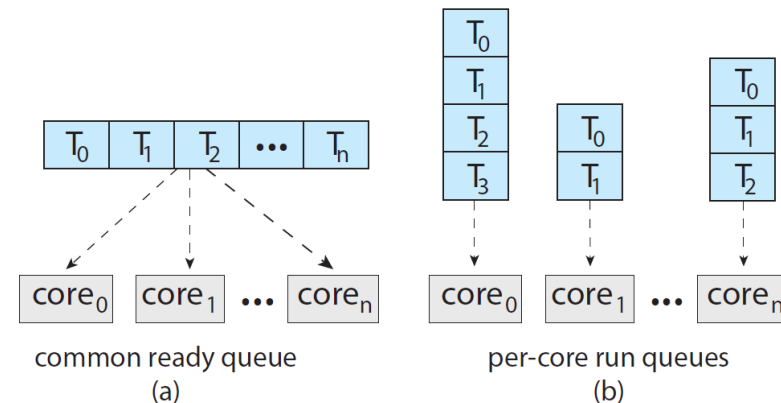
Min Quantum = 1
(1, 2, 4, 8, 16, 32)

Algorithm Evaluation

- How to select a CPU-scheduling algorithm for a system?
 - Based on **criteria/goal**
 - CPU utilization, response time, throughput, etc.
- Evaluation **methods**
 - Deterministic modeling (analytical evaluation)
 - Given algorithm and the (known) system workload, evaluate the performance
 - Queueing models (mathematical model)
 - Simulations (programming a model)
 - Actual-implementation (real-world testing)

Multicore/Multiprocessor Scheduling

- Multiple CPUs are available
 - Multicore CPUs
 - Multithreaded cores
 - NUMA systems
 - Heterogeneous multiprocessing
- **Load-sharing**
 - Processes/threads can run in parallel
- **Asymmetric** multiprocessing (AMP)
 - Common in embedded systems
- **Symmetric** multiprocessing (SMP)
 - **Widely** adopted (Linux, Windows, etc.)
 - ❑ All procs\thrs in **common** ready queue
 - ❑ Each processor **its own** ready queue
 - Work sharing/stealing
 - But, can specify **CPU affinity**



Linux Standard Scheduling, and API

- Different scheduling **classes**
 - Each class different policy (sort of MQ, but each “queue” different algorithm)
 - Implements POSIX policies
 - SCHED_FIFO, SCHED_RR – real-time (high prio)
 - SCHED_OTHER/SCHED_NORMAL, SCHED_BATCH – fair scheduling, **CFS algorithm**
 - SCHED_IDLE – idle scheduling
- To change **scheduling policy**
 - Process: `int sched_setscheduler(pid_t pid, int policy, const struct sched_param *param)`
 - Thread: `int pthread_attr_setschedpolicy(pthread_attr_t *tattr, int policy);`
- To change **scheduling priority**
 - Process: `int setpriority(int which, id_t who, int prio);`
 - Thread: `int pthread_setschedparam(pthread_t thread, int policy, const struct sched_param *param);`
- To change **CPU affinity** with *multiple CPUs* (non-POSIX)
 - `int sched_setaffinity(pid_t pid, size_t cpusetsize, cpu_set_t *mask);`

Summary

- Scheduling is a fundamental feature of OS
- Multiple goals, sometimes conflicting
- It can make a huge difference in performance
 - Difference increases with the variability in service requirements
- Many (single-CPU) algorithms
 - FCFS, SJF, RR, Priority, MQ, MLFQ
- Same algorithms adapted for multiple CPUs
- Evaluation of what algorithm is best is complex
- Real systems use hybrids