

Operating Systems
(INFR09047)
2019/2020 Semester 2

Memory Management:
Non-contiguous Allocations

abarbala@inf.ed.ac.uk

Chapter 9 (9.3, 9.4, 9.6)


Overview

- Non-contiguous memory allocation
- Segmentation
 - Basics
- Paging
 - Basics
 - Page Tables
 - TLB
 - Memory Protection
 - Shared Pages
 - Hierarchical Pages
 - Hashed Pages
 - Inverted Pages
 - Use cases

Memory Allocation

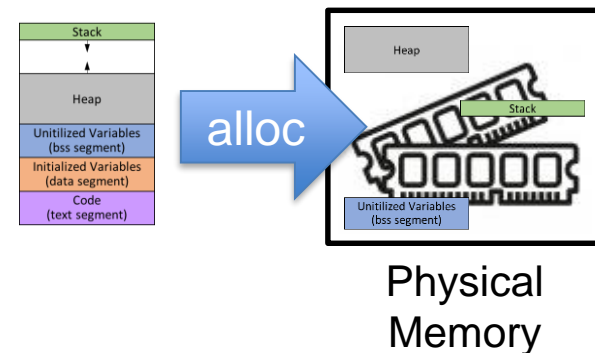
- **Contiguous**

- Allocation granularity is entire logical address space
- *Request physical space* for the **entire program at once**
 - **Contiguous** in logical memory and physical memory
- Issues
 - (external) Fragmentation
 - Long swap times
 - Long compaction times



- **Non-contiguous**

- Split logical address space in chunks
- *Request physical space* for each **program's chunk at the time**
 - **Contiguous** in logical memory, **non-contiguous** in physical
- Key mechanisms
 - Segmentation
 - Paging

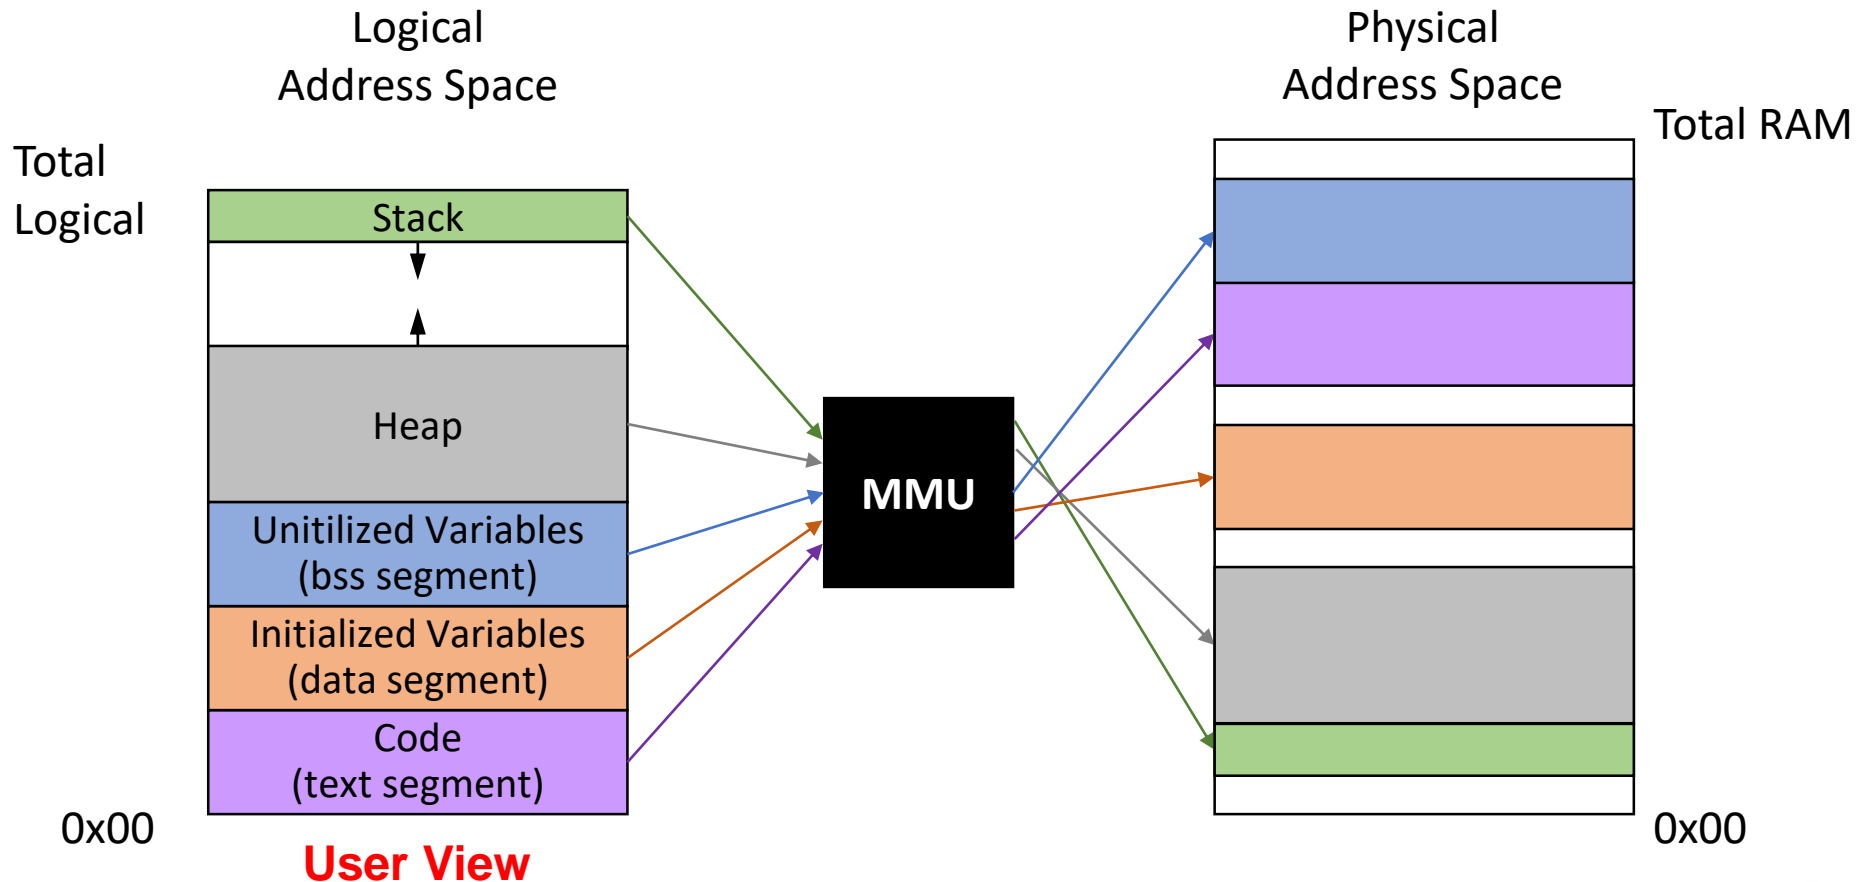


Segmentation

- Segmentation
 - Partition an address space into **variable size** chunks/units
 - *Logical units*
 - Stack, heap, data, code, subroutines, ...
 - A **logical address** is **<segment #, offset>**
- Facilitates sharing and reuse
 - Segment is a natural unit of sharing

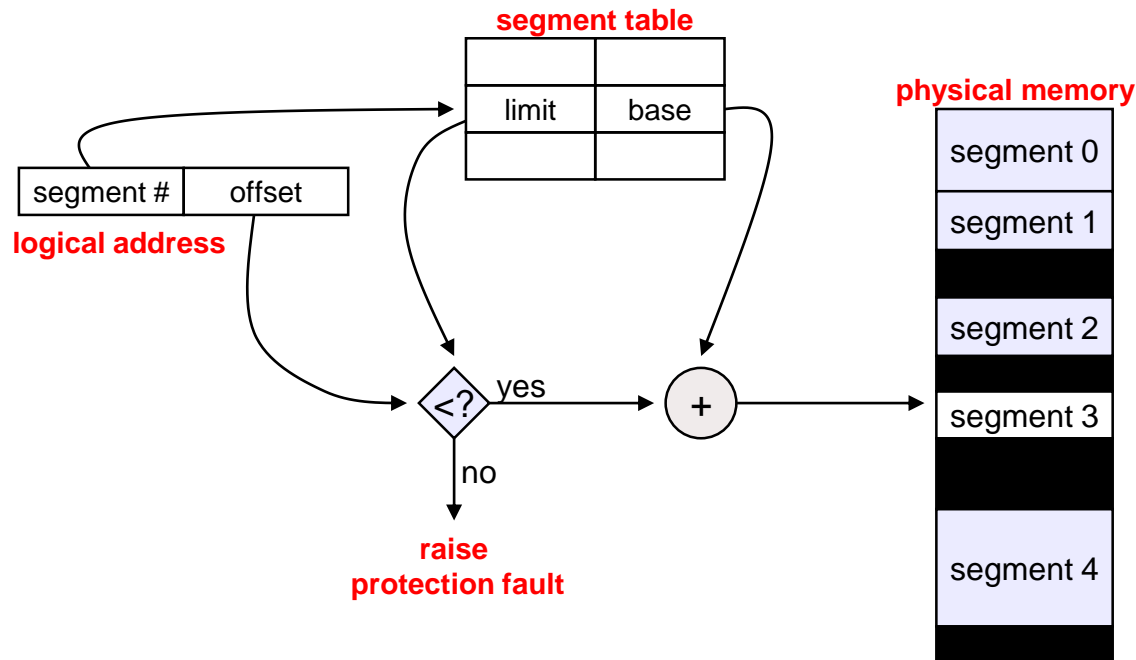
Segmentation

- *Logical address space* divided in **variable size** chunks



Hardware Support

- Segment table
 - Multiple base/limit pairs, **one per segment**
 - Segments named by **segment #**, used as **index into table**
 - A logical address is **<segment #, offset>**
 - Physical address = logical address offset + segment base address

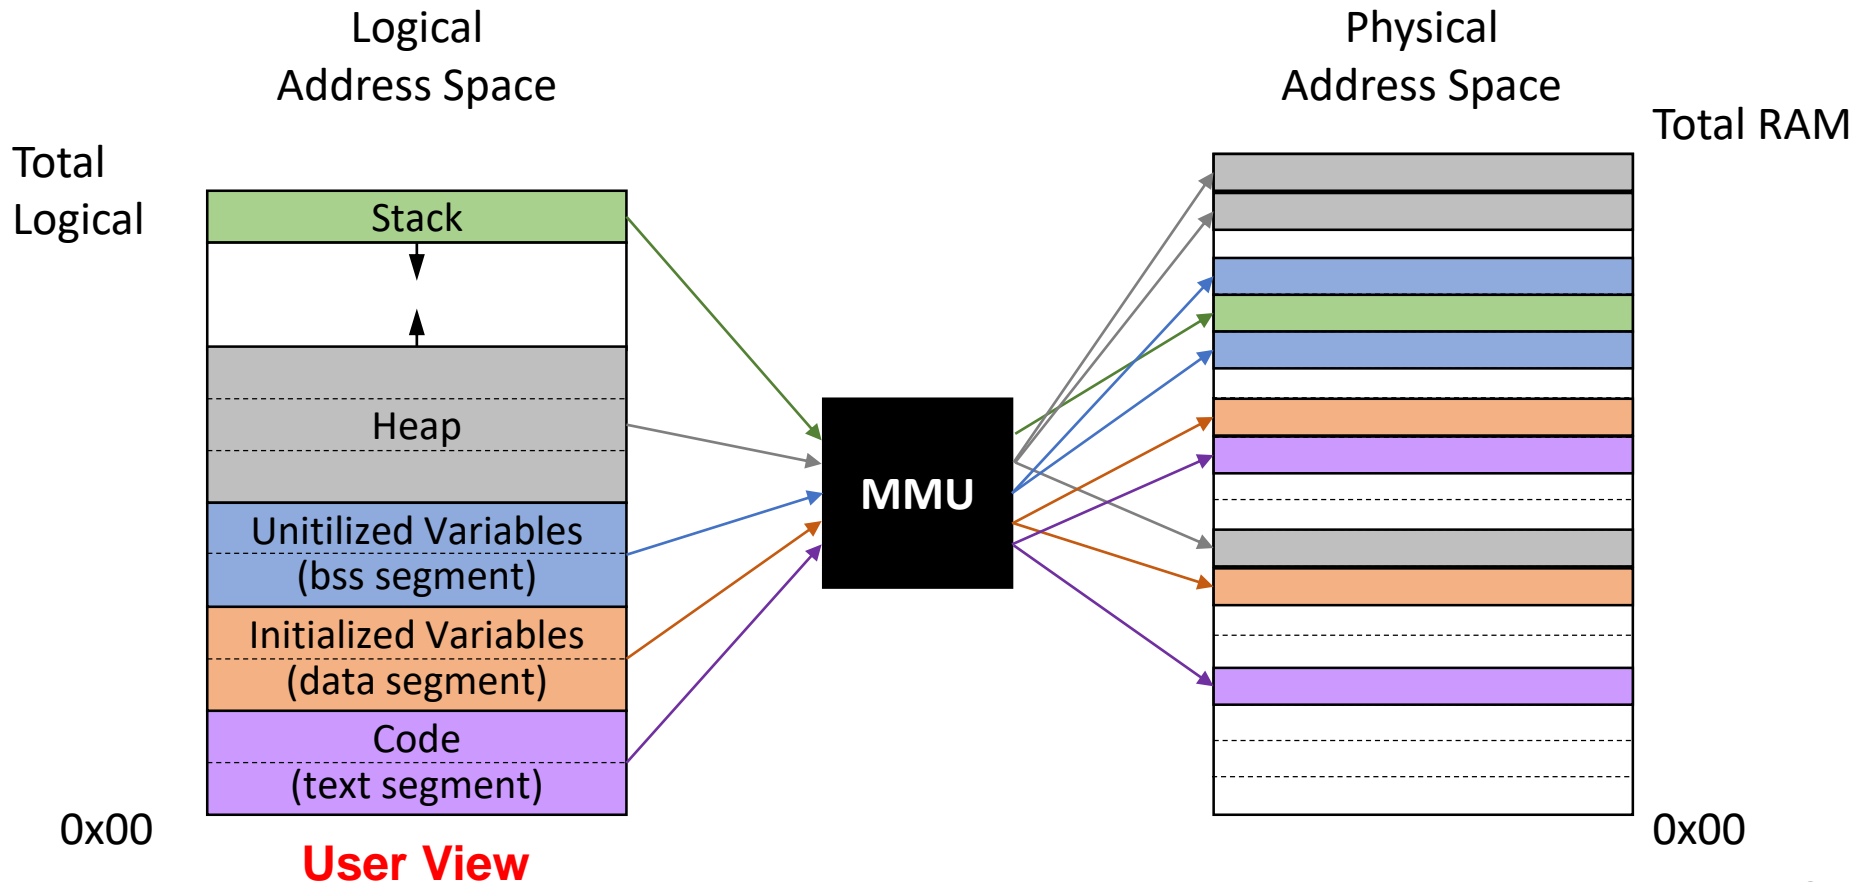


Pros and cons

- Allows non-contiguous physical addresses
 - Allocated “chunks” are smaller than entire program address space
 - Reduce fragmentation by exploiting varying sized holes
- Enables sharing
 - Same segment can be shared across processes
- Process view and physical memory **very different**
- By implementation, process can **only access** its own memory
- Segmentation rarely used today
 - x86 CPUs doesn't support it on 64bit models


Paging

- *Logical address space* divided in **fixed size** chunks
 - *Physical address space* divided in **fixed size** chunks, *same size*



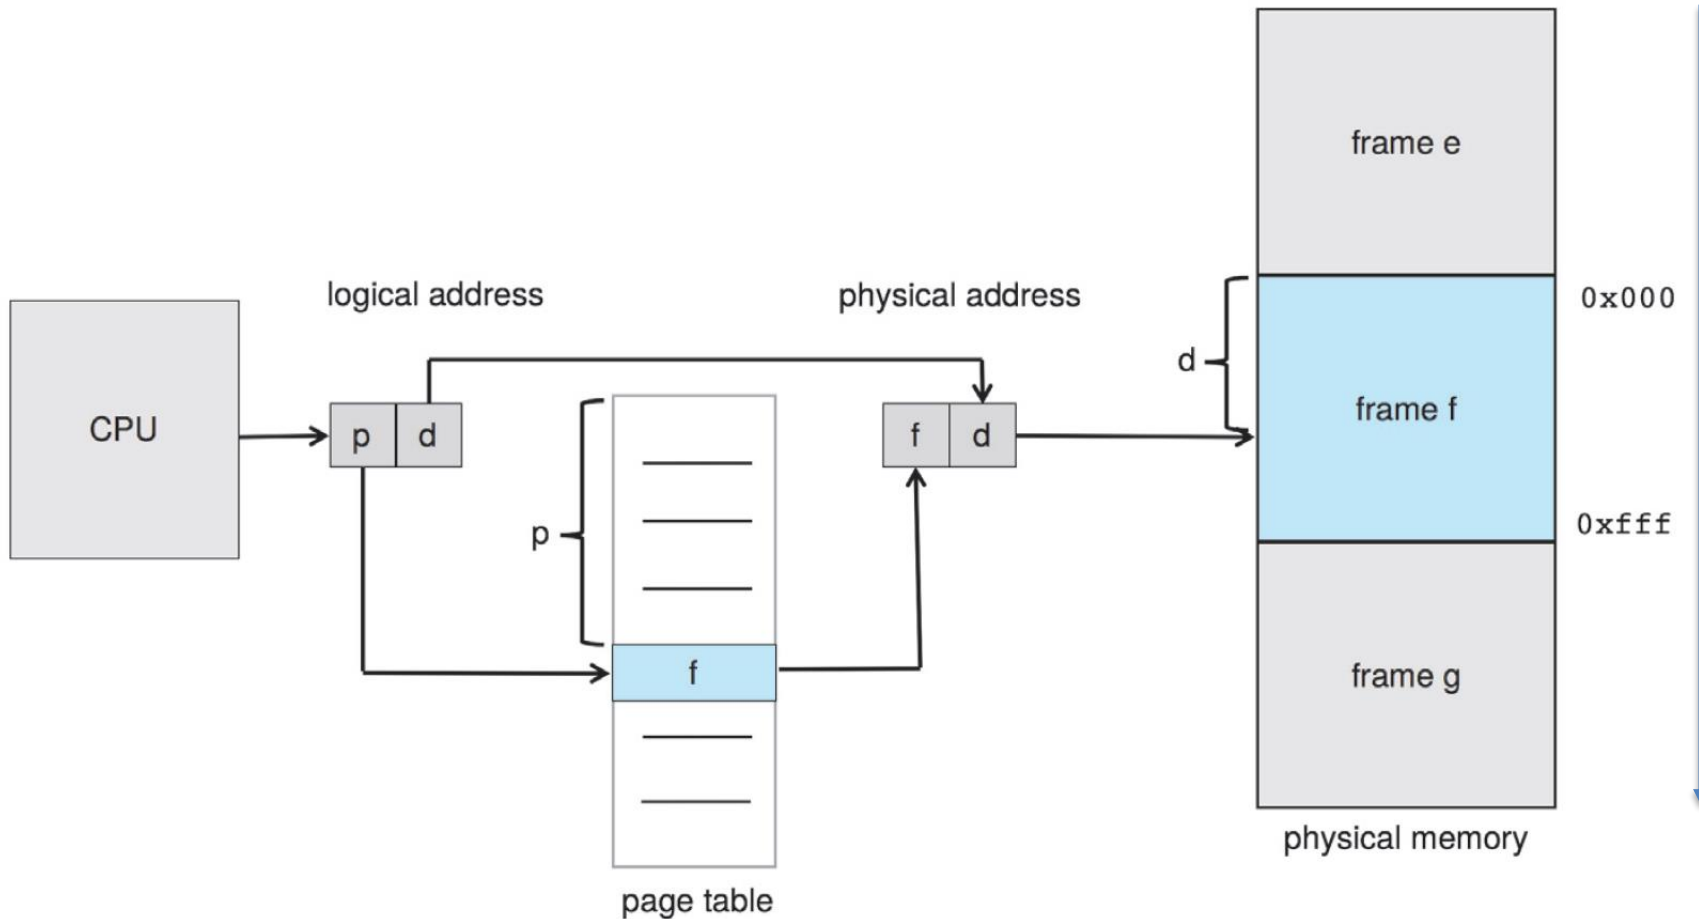
Address Translation Scheme

- Map **logical** chunks (pages) to **physical** chunks (frames)
 - **Page** size == **frame** size
 - Chunks are at a **predefined fixed** (logical and physical) address
- **Logical address** is divided into
 - **Page Number** (p)
 - Index into a **page table** which contains **frames base address**
 - **Page Offset** (d)
 - Summed to **frame base address** to make **physical** memory address



- Logical address space size **2^m bytes**
- Page and frame size **2^n bytes**

Paging Hardware




Page table: array of **page table entries**

Page table entry (PTE): frame base address and bits/flags

Paging Model of Logical and Physical Memory

**Page
Number**




logical
memory

0	1
1	4
2	3
3	7

page table

**Frame
Number**



physical
memory

Paging Example

Logical memory
addresses

Page
number

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

0

1

2

3

logical memory

0	5
1	6
2	1
3	2

page table

Physical memory
addresses

Frame
number

0	
4	i
	j
	k
8	m
	n
	o
	p
12	
16	
20	a
	b
	c
	d
24	e
	f
	g
	h
28	

0

1

2

3

4

5

6

7

physical memory

page number	page offset
p	d
m - n	n

$n=2$ and $m=4$


32-byte memory and 4-byte pages

Assume each character is 1-byte

Advantages with Paging


- **No external** fragmentation
- **Internal fragmentation depends on page size**
 - Page size = 2,048 bytes, process size = 72,766 bytes
 - 35 pages (71,680 bytes) + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - *Worst case* fragmentation = 1 frame – 1 byte
 - *On average* fragmentation = .5 frame size
 - *Are small frames desirable?*
 - Translations may require memory accesses (costly)
 - Different page sizes available on every system
- Process view and physical memory **very different**
- By implementation, process can **only access** its own memory

Free Frames



(a)

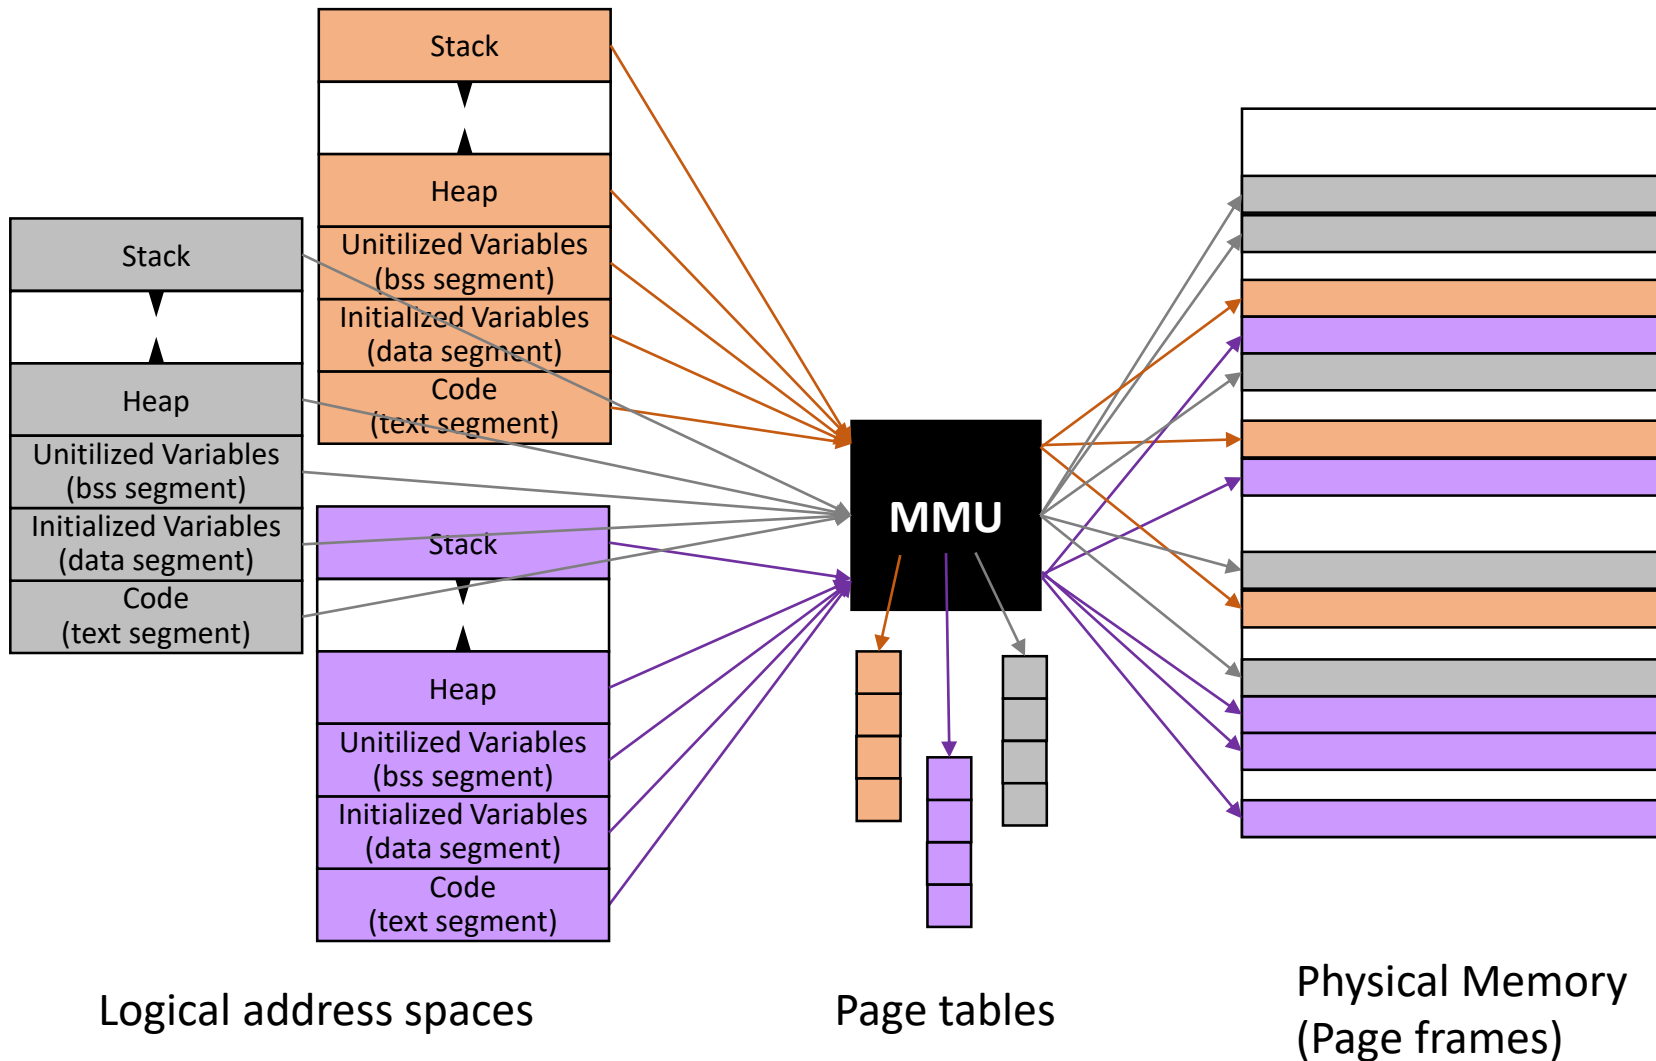
Before allocation



(b)

After allocation

Paging: Many Processes



Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)**
 - Points to the page table, **part of CPU**
- **Page-table length register (PTLR)**
 - Indicates size of the page table, **part of CPU**
- Every data/code access requires **two memory accesses**
 - To *fetch the **page table entry*** (translation)
 - To *fetch the actual **memory content*** (data/code)

Implementation of Page Table: TLB

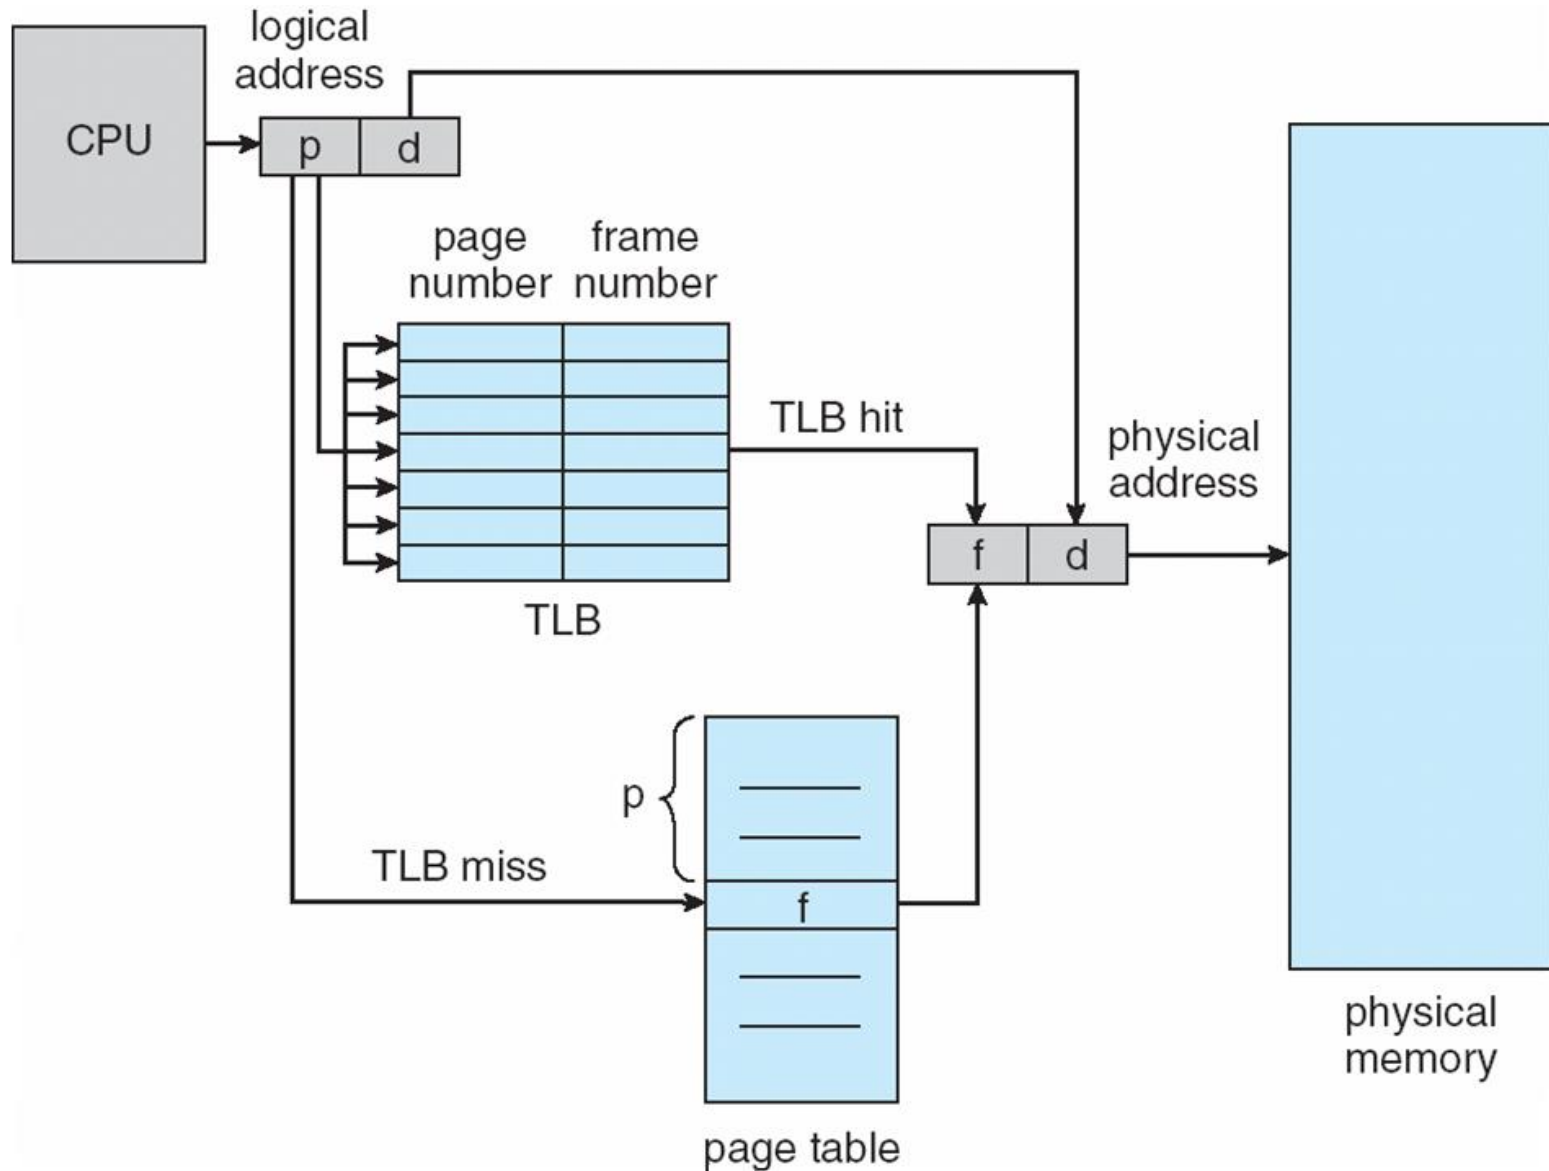
- How to avoid **two** memory accesses?
- **Translation look-aside buffers (TLBs)**

- Fast-lookup hardware cache
 - Associative memory for parallel search
 - Entry of the form **<page #, frame #>**
- Typically small (e.g., 64 or 1,024 entries)

Page #	Frame #
0	14
2	18

- If translation
 - **Exists** in TLB (hit), use it (at no additional cost)
 - **Doesn't exist** in TLB (miss), fetch translation from memory
- Maintain translations for subsequent memory accesses
 - **Replacement policies** must be considered
 - Some entries can be **wired down** for permanent fast access

Paging Hardware With TLB






Effective Access Time (EAT) with TLB

- Memory access time
 - Time the CPU waits to access main memory directly
- Hit ratio = α
 - Percentage of times that a page number is found in TLB
- Miss ratio = $1 - \alpha$
- Consider $\alpha = 80\%$, 100ns for memory access
 - $EAT = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider $\alpha = 99\%$, 100ns for memory access
 - $EAT = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

Memory Protection

- **Page Table Entries (PTEs)** contain more information
- Memory protection by **protection bits** in each PTE
 - What **type of access** is allowed?
 - Read-only, read-write, execute-only
 - Does a **translation exist**?
 - **Valid** indicates that the associated page
 - **is in the process' logical address space**, thus a legal page
 - **Invalid** indicates that the page
 - **is not in the process' logical address space**
 - ...
- Any violations result in a trap to the OS kernel

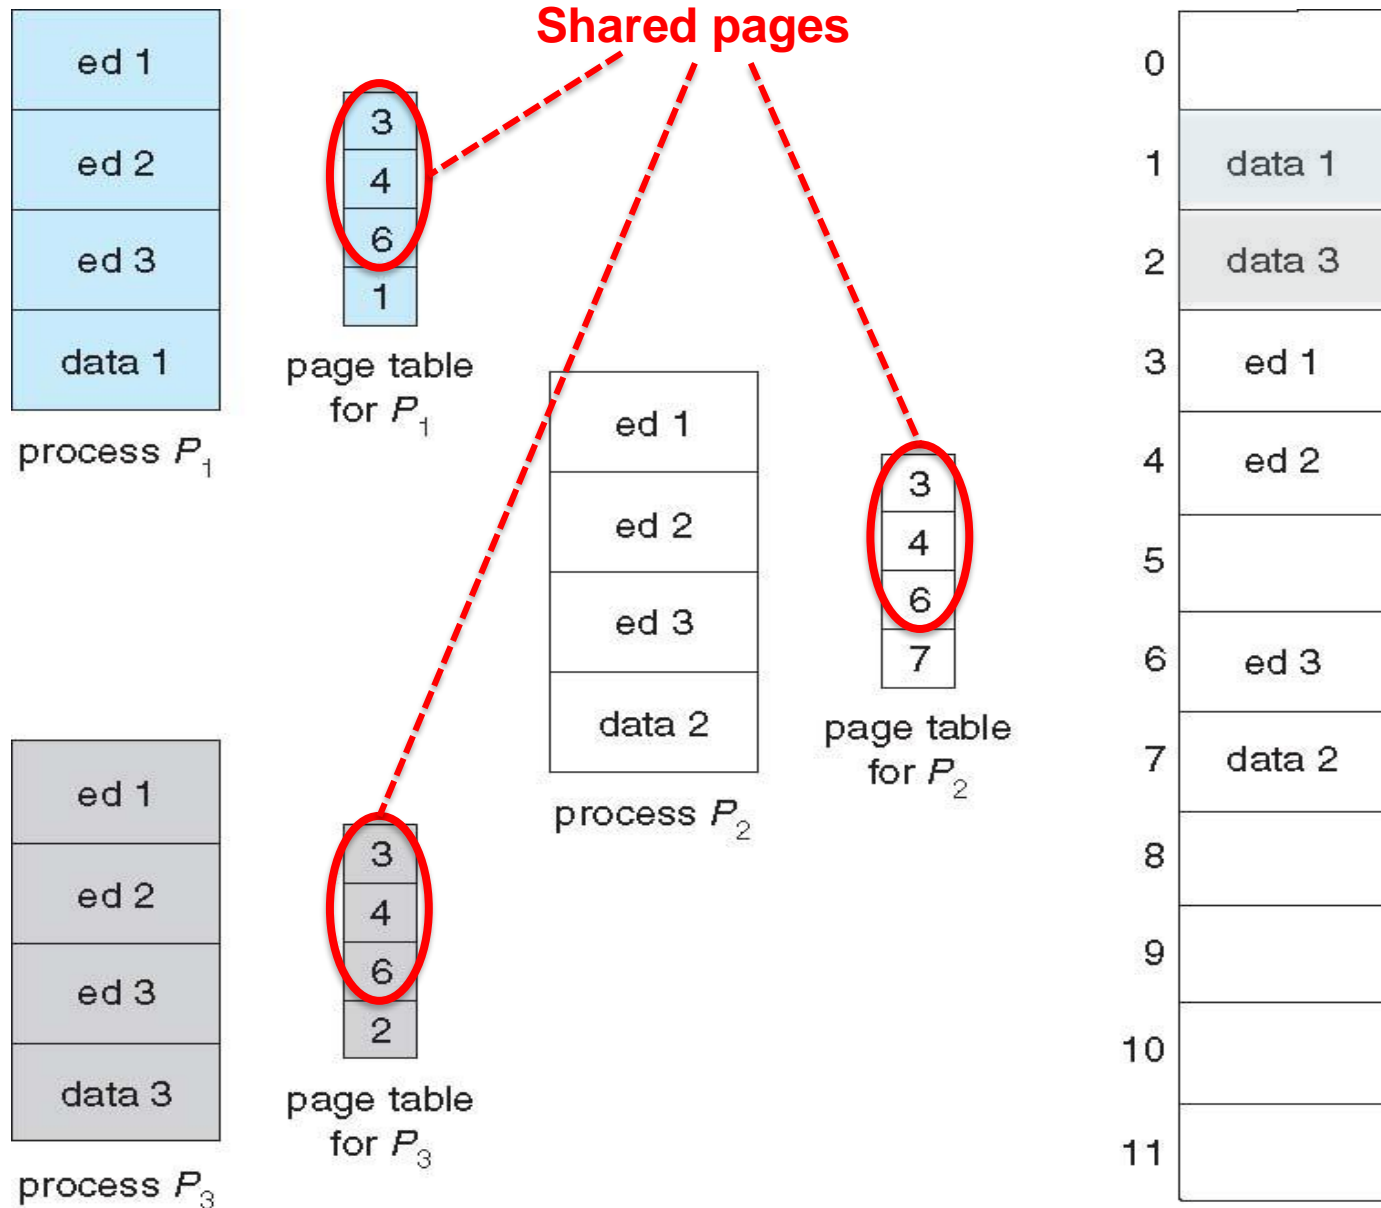
Memory Protection: **Valid (v)** or **Invalid (i)** Bits



Shared Pages

- **Shared code or data**
 - **One copy** of read-only (reentrant) **code** shared among processes
 - Different processes may share the same library code
 - But they will have per-process data
 - Read-write **data pages** shared among processes for **communication**
 - Inter-process communication
 - **Completely different from multithreaded processes**
 - Different threads share the entire address space, and OS resources
- **Private code and data**
 - Each process keeps a separate copy of the code and/or data
- Pages for the private/shared code and data can appear **anywhere** in the logical address space

Shared Pages Example




Structure of the Page Table

- Page table can **get huge**
 - Consider a 32-bit logical address space
 - Page size of 4 kB (2^{12})
 - Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - Each entry is 4 bytes, 4 MB of memory space required
 - Only for the page table
 - Need to allocate contiguously in physical memory
 - Costs a lot for small memories
 - 0.1% of the physical address space with 4GB
- Alternative **constructions**
 - **Hierarchical** Page Tables
 - **Hashed** Page Tables
 - **Inverted** Page Tables


Hierarchical Page Tables

- Break up the entire logical address space into **multiple** page tables
- Then construct **another table** that refers to each of such page tables
- **Two-level page table**



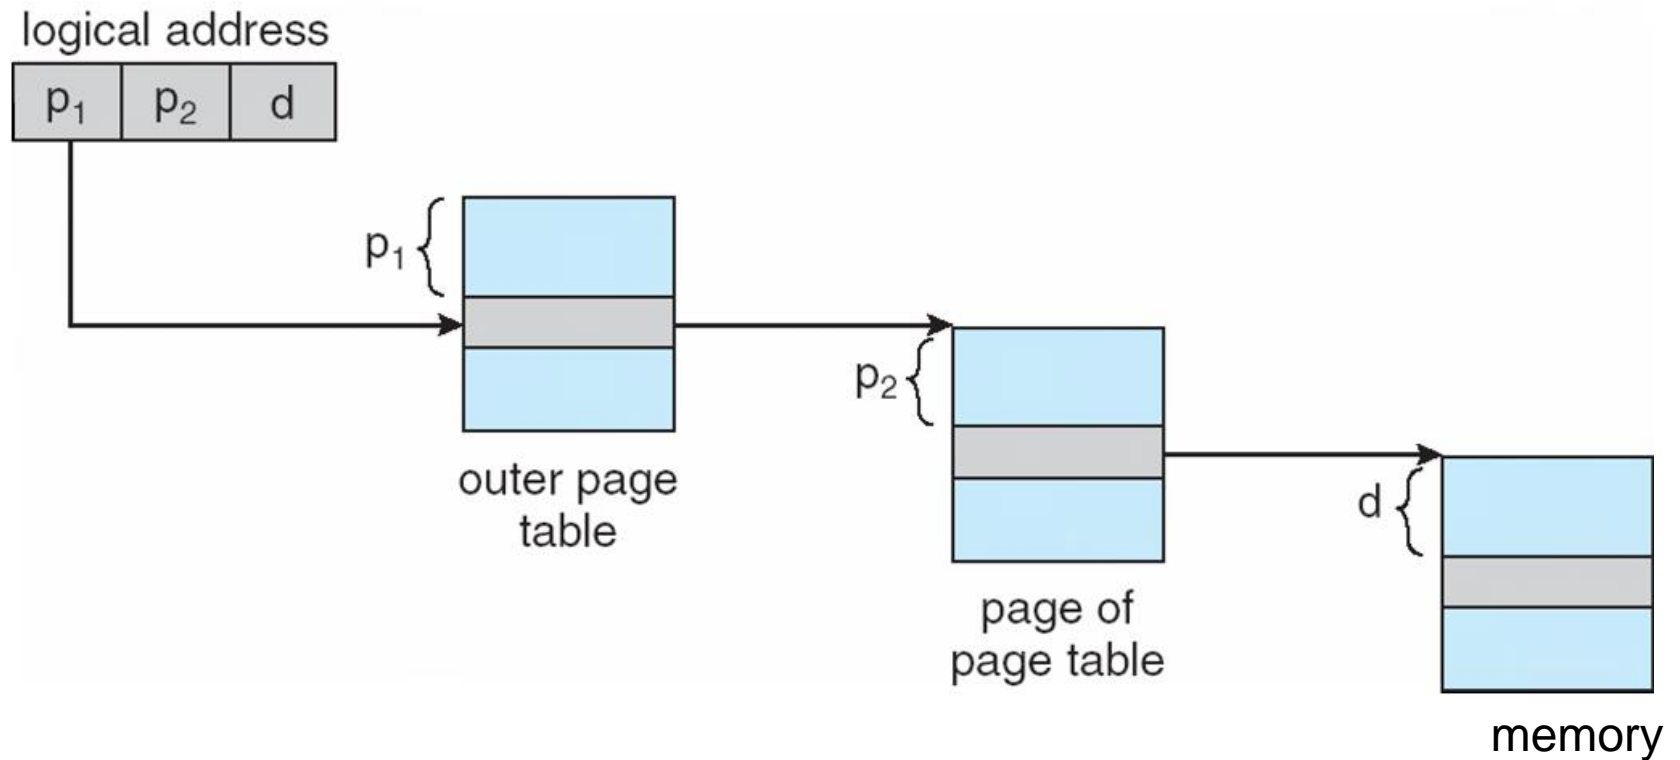
Two-Level Paging Example

- Logical address (32-bit machine with 1K page size) divided into
 - a **page number** consisting of 22 bits
 - a **page offset** consisting of 10 bits
- Page table is paged, the page number is further divided into
 - 12-bit **outer** page number
 - 10-bit **inner** page number
- A logical address looks like




- p_1 is an index into the outer page table
- p_2 is an index within the page of the inner page table

Two-Level Paging Address-Translation Scheme




64-bit Logical Address Space

- ❑ Two-level paging scheme not enough
- ❑ If page size is 4 kB (2^{12}) and **single-level** page table
 - ❑ Page table has 2^{52} entries
- ❑ If **two-level** scheme
 - ❑ Inner page tables could be 2^{10} entries
 - ❑ Outer page table has 2^{42} entries (4,398,046,511,104)



- ❑ If **three-level** scheme – add a 2nd outer page table
 - ❑ The 2nd outer page table is still 2^{32} entries (4,294,967,296)



Multi-level Paging Scheme

- On 64bit machines logical address space $< 2^{64}$
 - For practical reasons (no machine with 18 zetta bytes)
- But 4-level and 5-level page tables exist





Figure 2-1. Linear-Address Translation Using 5-Level Paging

Hashed Page Tables

- The **logical page number** is hashed into a page table index
- Each entry **chains** elements hashing to the same location
- Each element **contains**
 - Logical page number
 - Value of the mapped page frame
 - Pointer to the next element
- Logical page numbers are compared searching for a **match**
 - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
 - Similar to hashed but each entry refers to several pages
 - 16 rather than 1

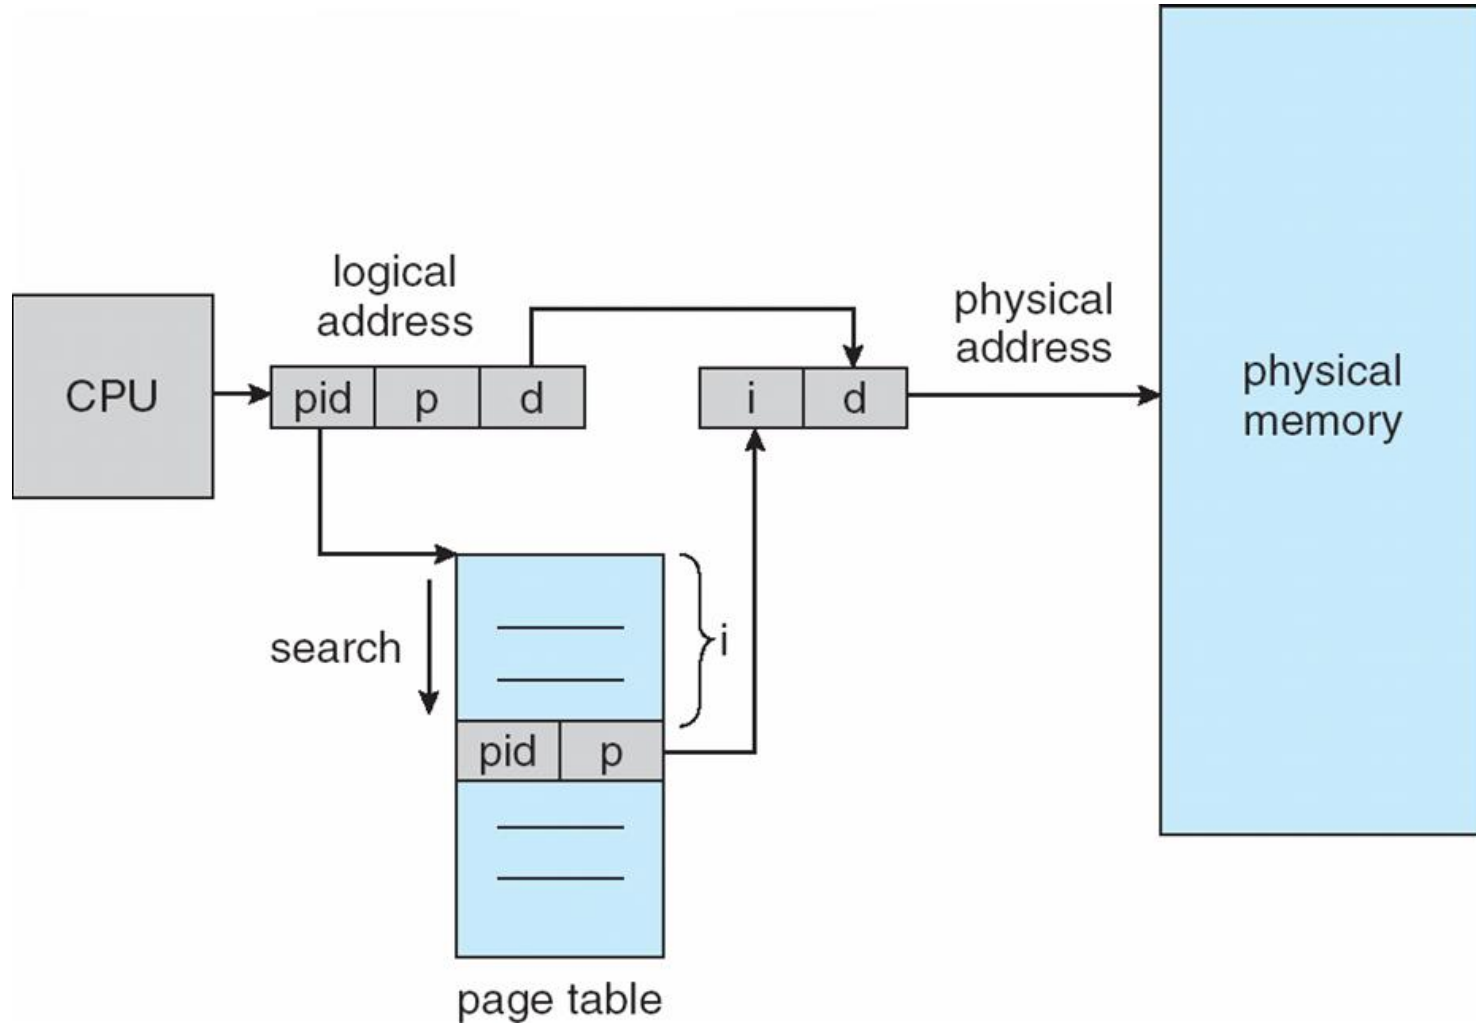
Hashed Page Table Example



Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages
 - **track all physical pages**
- One entry for each physical page of memory
- Entry consists of
 - **Virtual address** of the page stored in physical memory location
 - Information about the **process that owns** the page (protection)
- **Decreases memory** needed to store translations
 - Increases time needed to search the table
 - Use hash table to limit the search to one/few page-table entries
- How to implement shared memory?
 - Need OS intervention

Inverted Page Table Architecture



Why page tables? #1

- Inter-process memory protection
 - Process A address XYZ is different from Process B XYZ
- Protect code from being rewritten
 - Code pages read-only
 - A bad pointer can't change the program code
- Detect null pointer dereferencing at runtime
 - First page of the logical address space invalid
 - References to address 0 cause trap to OS
- Reduce memory usage (shared libraries)
 - Have one copy of a library in physical memory, not one per process
 - All page table entries to library point to the same set of physical frames
- Generalize use of “shared memory”
 - Regions of two processes' address spaces map to the same frames

Why page tables? #2

- **Copy-on-write (CoW)**, e.g. on fork()
 - Instead of copying all pages, create shared mappings of parent pages in child address space
 - Make shared mappings read-only for both processes
 - When either process writes, fault occurs, OS “splits” the page
- **Memory-mapped files**
 - Instead of using open, read, write, close
 - “map” a file into a region of the logical address space
 - e.g., into region with base ‘X’
 - Accessing logical address ‘X+N’ refers to offset ‘N’ in file
 - Initially, all pages in mapped region marked as invalid
 - OS reads a page from file whenever invalid page accessed
 - OS writes a page to file when evicted from physical memory
 - Only necessary if page is dirty

Summary

- Non-contiguous (physical) memory allocation
 - Reduced or no fragmentation
 - Enable sharing
 - Require hardware support
- Segmentation
 - Variable size chunks
- Paging
 - Fixed size chunks
 - Multiple optimizations