


Operating Systems
(INFR09047)
2019/2020 Semester 2

Synchronization

abarbala@inf.ed.ac.uk

Why Synchronization?


- **Cooperating tasks** access the same data
 - Two or more **threads** of the same process
 - Shared *data* and code
 - Two or more **processes**
 - *Data* on shared memory
 - ~~*Data* exchange via message passing~~



- **Concurrent or parallel** access to shared data
 - **Concurrent** processes/threads *virtually* run *at the same time* on a core
 - Switch between them at any time – *scheduling*
 - **Parallel** process/threads run *at the same time* on different cores
- May result in **data inconsistency** (integrity of data)
 - How to ensure **ordered** execution of cooperating tasks?
 - Synchronization

Ordered Execution: Temporal Relations

- Instructions executed by a single thread/process are **totally ordered**
 - $A < B < C < \dots$
 - $A' < B' < \dots$
 - $X < Y$ means X event **happened before** Y event
- In **absence of synchronization**
 - Instructions executed by distinct threads/processes must be considered **unordered / simultaneous**
 - On **one** core, or **multiple** cores
 - Not $X < X'$, and not $X' < X$
 - $C == A'$ or $C == B'$
- **Creation** relations always hold
 - But $A < B < A' < B' < \dots$



Example: Shared Bank Account #1

- A function to **withdraw money** from a bank account

```
int withdraw(account, amount) {  
    int balance = get_balance(account);           // read  
    balance -= amount;                             // modify  
    put_balance(account, balance);                 // write  
    return balance  
}
```

- You and your mother **share a bank account** with a balance of £1500.00
- What happens if you both go to **separate** ATM machines, and **simultaneously** withdraw £50.00 from the account?



Example: Shared Bank Account #2

- Bank's application is **multi-threaded**
- A separate **thread for each ATM** doing a withdrawal
 - Both threads run on the same bank server
- Each thread can **context switch** after each instruction

Thread YOU

```
int withdraw(account, amount) {  
    int balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    return balance;  
}
```



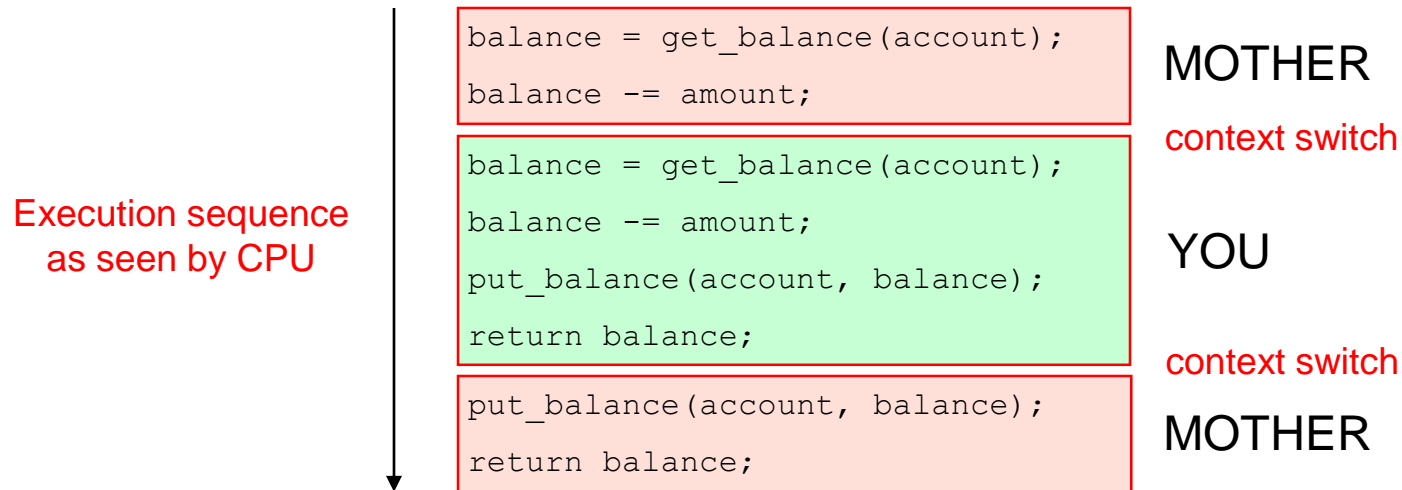
Thread MOTHER

```
int withdraw(account, amount) {  
    int balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    return balance;  
}
```



Example: Shared Bank Account #3

- **Interleaved** execution of the two threads



- What is the account balance after this sequence?
 - Who is happy, the bank or you and your mother?
 - How often is this sequence likely to occur?

Example: Shared Bank Account #4

- Which **interleavings** are OK?
- Which are not?

Thread YOU

```
int withdraw(account, amount) {  
    int balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    return balance;  
}
```



Thread MOTHER

```
int withdraw(account, amount) {  
    int balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    return balance;  
}
```



Example: Producer-consumer Problem #1

- Also known as the bounded-buffer problem
 - Fixed-size buffer ***B*** (with ***n*** elements)
 - ***p*** producer processes
 - ***c*** consumer processes
 - Producer and consumer processes share the buffer **B**
 - Producer process puts info into the buffer **B**
 - Consumer process takes info out of the buffer **B**



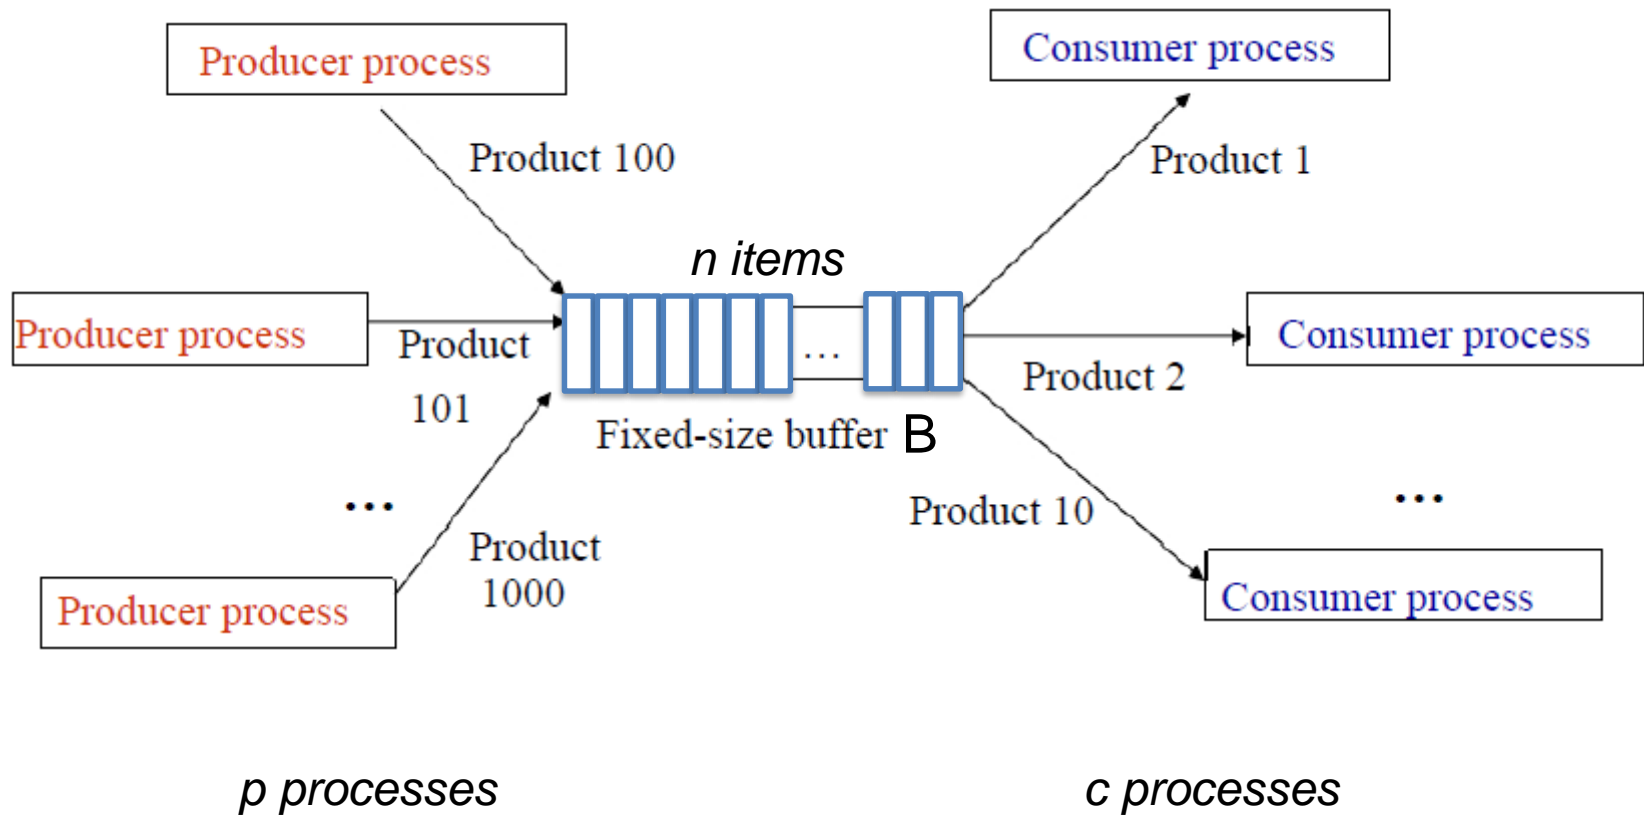
Barber shop analogy

n = number of waiting seats

p = all clients

c = barbers

Example: Producer-consumer Problem #2



Example: Producer-consumer Problem #3

- **Single-producer, single-consumer** problem ($p=1, c=1$)

Shared variable:
const int n;
int count=0;
Item buffer [n];

Producer

```
while (1){  
    ...  
    produce an item A;  
    ...  
    while(count==n);  
    insert(item);  
    count++;  
}
```

Consumer

```
while (1){  
    while(count==0);  
    item=remove_item();  
    count--;  
    ...  
    consume an item;  
}
```

Example: Producer-consumer Problem #4

- `count++`; `count--`;
instructions are **not guaranteed** to execute as a single machine instruction

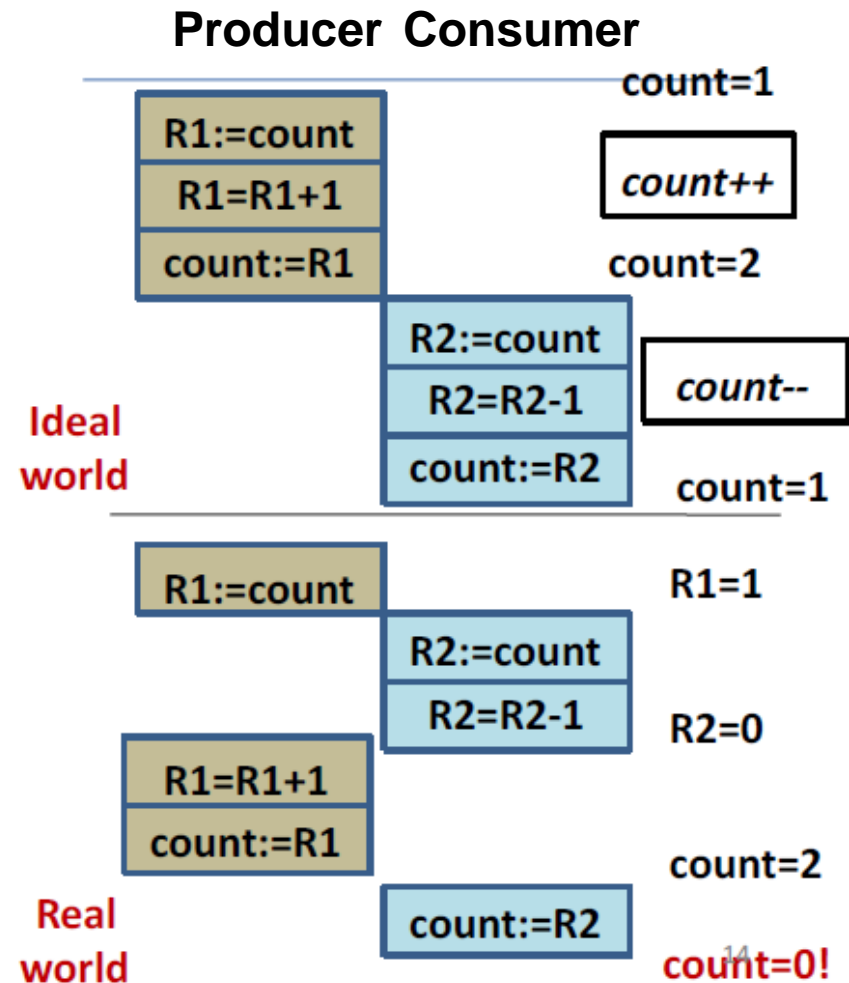
- `count++`; a possible assembly implementation

```
R1 := count;
```

```
R1 = R1+1;
```

```
count := R1;
```

Note R1 is a machine register




Race Conditions

- Examples
 - Results of concurrent or parallel access is **non-deterministic**
 - **Depends on**
 - Timing
 - When context switches occurred
 - What process/thread ran at context switch
- **Race condition**
 - Two or more processes reading or writing shared data
 - The final result depends on who runs precisely when
- How **to avoid** race conditions?

Modeling Programs to Solve Race Conditions

- **Critical Region or Section**
 - Part of the program where the **shared data is accessed**
 - **Uncoordinated** read/write of the data in critical section may lead to races



- **Common pattern to identify**
 - Read-modify-write of a shared data (variable)
 - Globals and heap-allocated variables
 - In code that can be executed by concurrent or parallel threads

Critical Region in Bank Example

Thread YOU

```
int withdraw(account, amount) {  
    int balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    return balance;  
}
```

Thread MOTHER

```
int withdraw(account, amount) {  
    int balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    return balance;  
}
```

From read to write the balance

Critical Region in Producer-consumer Example

Access to the count variable

Shared variable:
const int n;
int count=0;
Item buffer [n];

Producer

```
while (1){  
    ...  
    produce an item A;  
    ...  
    while(count==n);  
    insert(item);  
    count++;  
}
```

Consumer

```
while (1){  
    while(count==0);  
    item=remove item();  
    count--;  
    ...  
    consume an item;  
}
```

Avoid Race Conditions

- Find some way to **prohibit more than one process** in its critical region(s) at the same time
- Solution: **Mutual exclusive** access to critical regions
 - Some way of making sure that
 - If one process is using a shared variable
 - The other processes will be excluded from doing the same thing
 - **Only one** process/thread in a critical section **at any time**

Mutual Exclusion



This is not an example of mutual exclusion

Requirements to Avoid Race Conditions

- Critical Regions are **not enough**
 - **Mutual exclusion**
 - At most one thread/process is in the critical section
 - **Progress**
 - No process running outside its critical region may block other processes
 - **Bounded waiting**
 - No process should have to wait forever to enter its critical region
 - **Performance**
 - Overhead of entering and exiting critical section is small wrt the work being done within it
 - No assumption can be made about processing speed

How the Code of a **Critical Region** Looks Like?

- Pseudocode example

repeat

do other work */* do other work */*

enter_critical */* enter critical region, may wait to enter */*

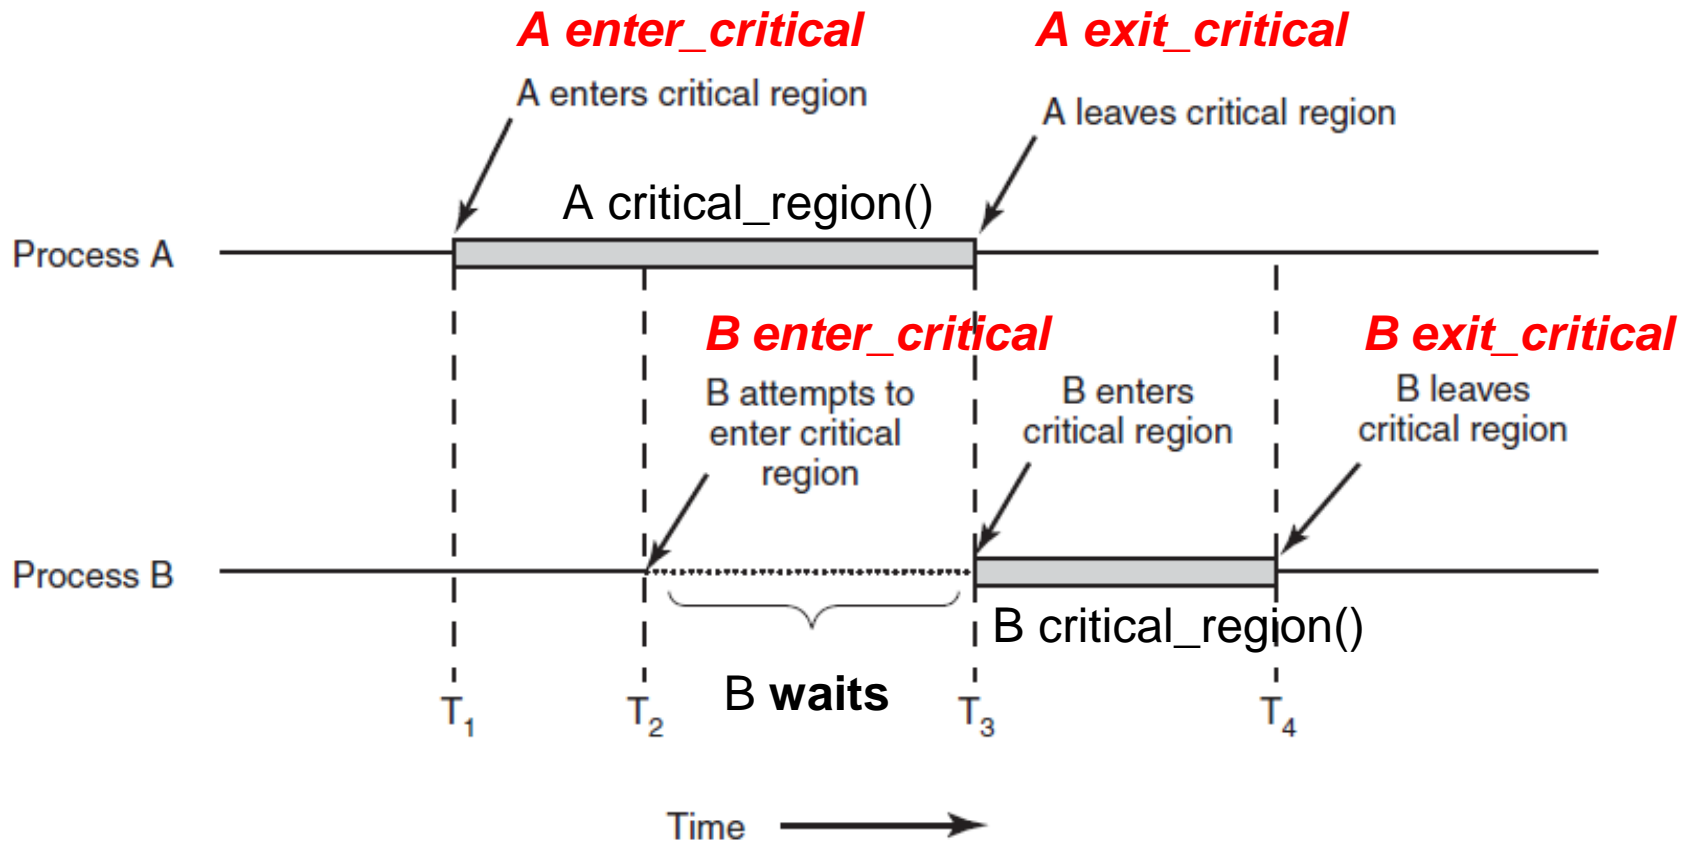
critical_region() */* access shared variables */*

exit_critical */* leave critical region */*

remainder_region */* do other work */*

until *condition*

- How to implement *enter_critical*, and *exit_critical* to **guarantee mutual exclusion?**



Mechanisms


- **Disable Interrupts*****
 - Needs operating systems, high overhead (especially on multi-CPU)
- **Locks/Spinning locks (Spinlocks)**
 - Primitive, minimal semantics, used to build others
- **Semaphores (and non-spinning locks)**
 - Basic, easy to get the hang of, somewhat hard to program with
- **Monitors**
 - Higher level, require language support, easier to program with

Locks


- A lock is a an object with methods
 - **acquire()**: obtain the right to enter the critical section
 - **Prevents progress** of the task until the lock is acquired
 - **release()**: give up the right to be in the critical section
 - Immediate
- Note
 - Terminology varies
 - `acquire()` / `release()`
 - `lock()` / `unlock()`

Locks: Example


→ is the "happens-before" relation



Possibly incorrect




Correct



Correct

How to
implement
with
locks?



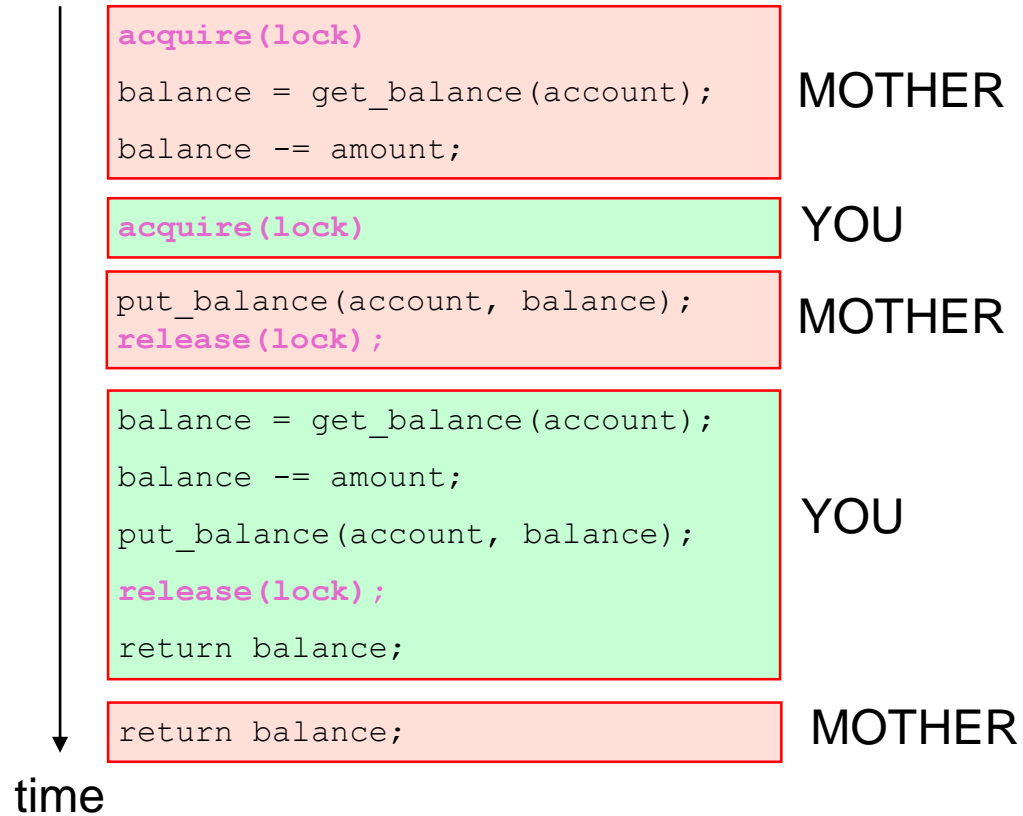
Acquire/Release

- Programmer **pairs up** calls to `acquire()` and `release()`
 - Per thread
 - Or single threaded process
 - From `acquire()` to `release()`
 - A thread **holds the lock**
 - `acquire()` does not return until the caller “owns” (holds) the lock
 - Thread waits (blocked) to enter its critical section
 - At most **one thread** can hold a lock at any time
- What happens if calls aren't paired?
- What happens if two threads acquire different locks?
- What is the right granularity of locking?

Using Locks

```
int withdraw(account, amount) {  
    acquire(lock) ;  
    balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    release(lock) ;  
    return balance;  
}
```

} critical
section



How to Implement Locks?

- Spinning on a lock variable

```
struct lock_t {  
    int held = 0;  
}  
  
void acquire(lock_t* lock) {  
    while (lock->held) { };  
    lock->held = 1;  
}  
  
void release(lock_t* lock) {  
    lock->held = 0;  
}
```

```
while (TRUE) {  
    acquire(lock);  
    critical_region(); /* work */  
    release(lock);  
    noncritical_region();  
}
```

spinning
}

```
1 while (TRUE) {  
2     while (lock->held != 0); /* loop */  
3     lock->held = 1;  
4     critical_region(); /* work */  
5     lock->held = 0;  
6     noncritical_region();  
7 }
```

T1

```
1 while (TRUE) {  
2     while (lock->held != 0); /* loop */  
3     lock->held = 1;  
4     critical_region(); /* work */  
5     lock->held = 0;  
6     noncritical_region();  
7 }
```

T2

Spinning on a Lock Variable

- **Race condition** may happen
 1. Process T1 sees lock=0 moves to line 3
 2. Process T1 is descheduled and Process T2 is scheduled
 3. Process T2 sees lock=0, moves to line 3, sets lock =1, **enters critical section**
 4. Process T2 is descheduled and Process T1 is scheduled
 5. Process T1 sets lock=1, **enters critical section**

```
1 while (TRUE) {
```

```
2 ① while (lock->held != 0); ② /* loop */
```

```
3  lock->held = 1; ④
```

```
4 ⑤ critical_region(); /* work */
```

```
5  lock->held = 0;
```

```
6  noncritical_region();
```

```
7 }
```

T1

```
1 while (TRUE) {
```

```
2 ② while (lock->held != 0); ② /* loop */
```

```
3  lock->held = 1; ③
```

```
4  critical_region(); ④ /* work */
```

```
5  lock->held = 0;
```

```
6  noncritical_region();
```

```
7 }
```

T2

How to Fix That?

- **Problem**
 - Implementation of spinning on a lock variable **has critical sections**
 - The acquire/release must be **atomic**
 - Executes “**all or nothing**”
 - From the view of all CPUs/cores
- Can we solve with **Software**?
 - **Strict alternation**
 - “Progress” property missing
 - **Peterson Solution**
 - Doesn’t easily scale to multiple threads
 - Locking in the OS
 - It works only for user code, very expensive
- Need **hardware**
 - Atomic instructions (Test-and-set, compare-and-swap, ...)

Hardware Test-and-Set

- CPU provides the following as a **single atomic instruction**

```
bool test_and_set(bool *flag) {  
    bool old = *flag;  
    *flag = True;  
    return old;  
}
```

- Different CPUs implement it differently
 - But in **one assembly instruction**
 - TAS register, flag_address

Implementing Locks with Test-and-Set

- Spinning on a lock variable (with TAS)

```
struct lock_t {
    int held = 0;
}

void acquire(lock_t* lock) {
    while(test_and_set(&lock->held)) { };
}

void release(lock_t* lock) {
    lock->held = 0;
}
```

```
while (TRUE) {
    acquire(lock);
    critical_region(); /* work */
    release(lock);
    noncritical_region();
}
```

```
1 while (TRUE) {
2     acquire:
3     TAS REGISTER, &(lock->held)
4     CMP REGISTER, #0
5     JNE acquire
6     critical_region(); /* work */
7     release:
8     MOV &(lock->held), #0
9     noncritical_region();
10}
```

T1

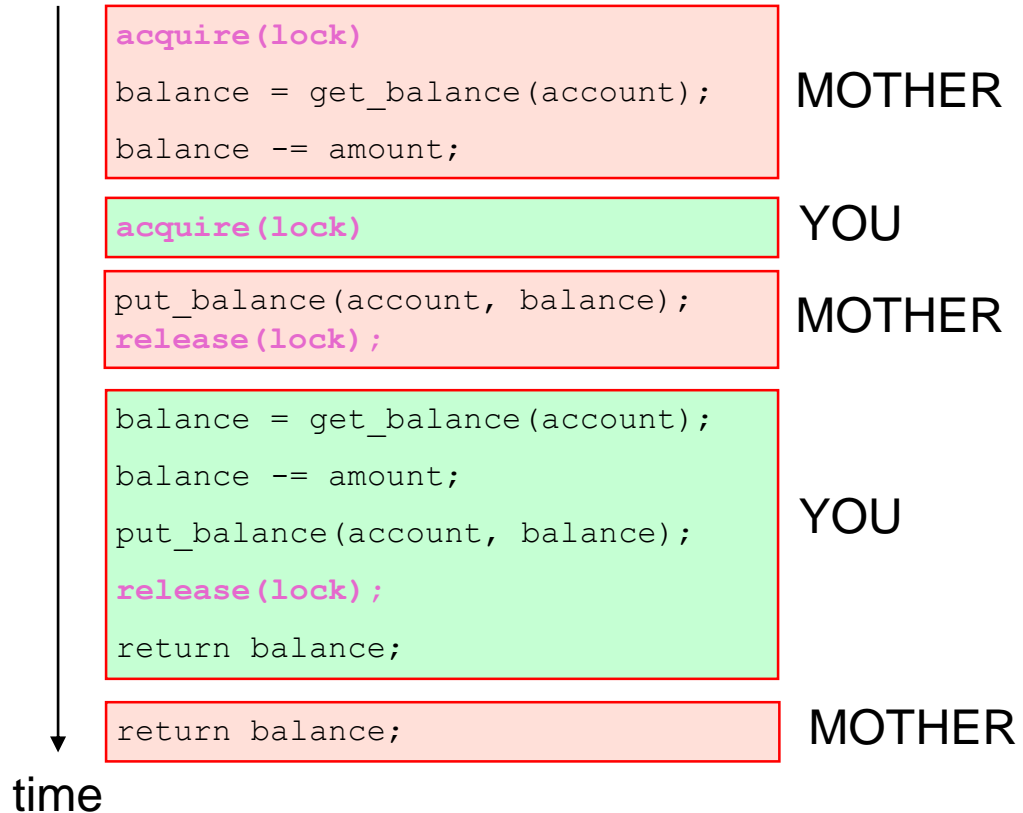
```
1 while (TRUE) {
2     acquire:
3     TAS REGISTER, &(lock->held)
4     CMP REGISTER, #0
5     JNE acquire
6     critical_region(); /* work */
7     release:
8     MOV &(lock->held), #0
9     noncritical_region();
10}
```

T2

Using Locks with Test-and-Set

```
int withdraw(account, amount) {  
    acquire(lock);  
    balance = get_balance(account);  
    balance -= amount;  
    put_balance(account, balance);  
    release(lock);  
    return balance;  
}
```

} critical section



- How does a thread blocked on an “acquire” (that is, stuck in a test-and-set loop) **yield** the CPU?
 - **Voluntarily** calls yield() (*spin-then-block*)
 - **Involuntary** context switch (e.g., timer interrupt)

Spinning on Lock Variables ...

- Spinlocks
 - **Wastes CPU resources**
 - When a thread is spinning on a lock
 - Thread holding the lock **cannot make any progress**
 - **Just burns** CPU cycles
- How to solve this?
- Spinlocks
 - **Primitives** to build higher-level synchronization constructs
 - Ensure acquiring only happens for a short time

Summary

- Threads/processes need access shared data
- Synchronization
 - Introduces temporal ordering
 - May eliminate races
 - Provided by
 - Disable interrupts***
 - Locks (this class)
 - Semaphores (next class)
 - Monitors (next class)
- Spinlocks are the lowest-level mechanism
 - Primitive in terms of semantics
 - Error-prone!
 - Implemented by spin-waiting
 - Crude!