

Operating Systems
(INFR09047)
2019/2020 Semester 2

Semaphores and Monitors

abarbala@inf.ed.ac.uk

What About Busy Waiting Solutions?

- Disadvantages

- Waste CPU time waiting
- High priority jobs may hinder progress of low priority jobs (scheduling)
- Protect a single resource

```
1 while (TRUE) {  
2   enter_region:  
3     TAS REGISTER, LOCK_ADDR  
4     CMP REGISTER, #0  
5     JNE enter_region  
6   critical_region(); /* work */  
7   leave_region:  
8     MOV LOCK_ADDR, #0  
9   noncritical_region();  
10 }
```

Process A and Process B

- a) Process A enters its critical region
- b) Process A descheduled, Process B scheduled (high priority)
- c) Process B busy waits for the critical region (depends on A)
- d) Process B will busy wait until Process A is scheduled back
- e) A is scheduled back and exits its critical section

sleep() and wakeup()

- Instead of busy wait, let the process sleep
- Introduce new **Operating System** functions
- `sleep()`
 - Caller gives up the CPU for some duration of time
 - Until a thread/process wakes it up
- `wakeup()`
 - Caller wakes up some sleeping thread/process
- In practice
 - As `sleep()` a thread/process may call `yield()` syscall
 - A thread/process may be woken with a signal
 - Each has also a kernel-level implementation

Semaphore

- E.W. Dijkstra, 1965
 - Semaphore variable ***value***

```
int value=N; /* protected resource */
```

function ENTER_CRITICAL

```
– Proberen (value) /* P, Wait */
```

```
value--;
```

```
if (value < 0)
```

```
    sleep(); /* sleep without completing wait */
```

function EXIT_CRITICAL

```
– Verhogen (value) /* V, Signal */
```

```
value++;
```

```
if (value <= 0)
```

```
    wakeup(); /* wakeup one sleeper with a policy */
```

Each function executes atomically

Semaphore Implementation

```
typedef struct {  
    int value;  
    struct thread *list;  
} semaphore;  
  
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this thread to S->list;  
        sleep();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a thread T from S->list;  
        wakeup(T);  
    }  
}
```

- A list of TCBs
- The list kept in any **order** (e.g. FIFO)
- User- or kernel-level TCBs
- Each function should be atomic

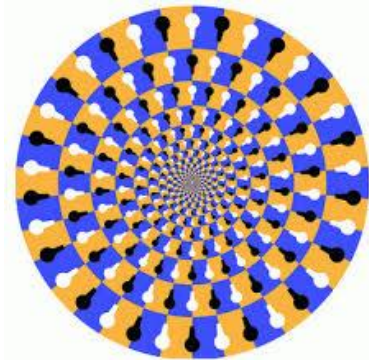
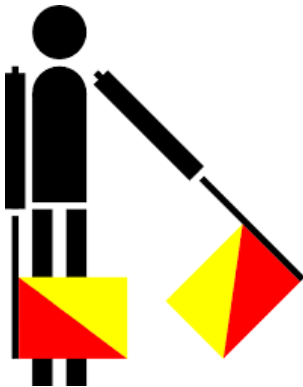
How to Provide Atomicity?

Need a spinlock! (“real” mutual exclusion lock)

- P/wait(sem)
 - acquire “real” mutual exclusion lock
 - if sem is “available” (>0), decrement sem; release “real” mutual exclusion lock; let thread continue
 - otherwise, place thread on associated queue; release “real” mutual exclusion lock; run some other thread
- V/signal(sem)
 - acquire “real” mutual exclusion lock
 - if thread(s) are waiting on the associated queue, unblock one (place it on the ready queue)
 - if no threads are on the queue, sem is incremented
 - » the signal is “remembered” for next time P(sem) is called
 - release “real” mutual exclusion lock
 - the “V-ing” thread continues execution

Semaphores vs Spinlocks

- **With Semaphores**
- Threads/processes are blocked at the level of **program logic**
 - By the semaphore P/wait operation
 - Placed on queues, rather than busy-waiting
 - The scheduler is **aware** about thread/process **waiting**
 - **Other tasks** can execute
- Busy-waiting may be used for the “**real**” **mutual exclusion lock**
 - To implement P/wait and V/signal
 - These are very short critical sections – **independent of program logic**
 - They are not implemented by the application programmer

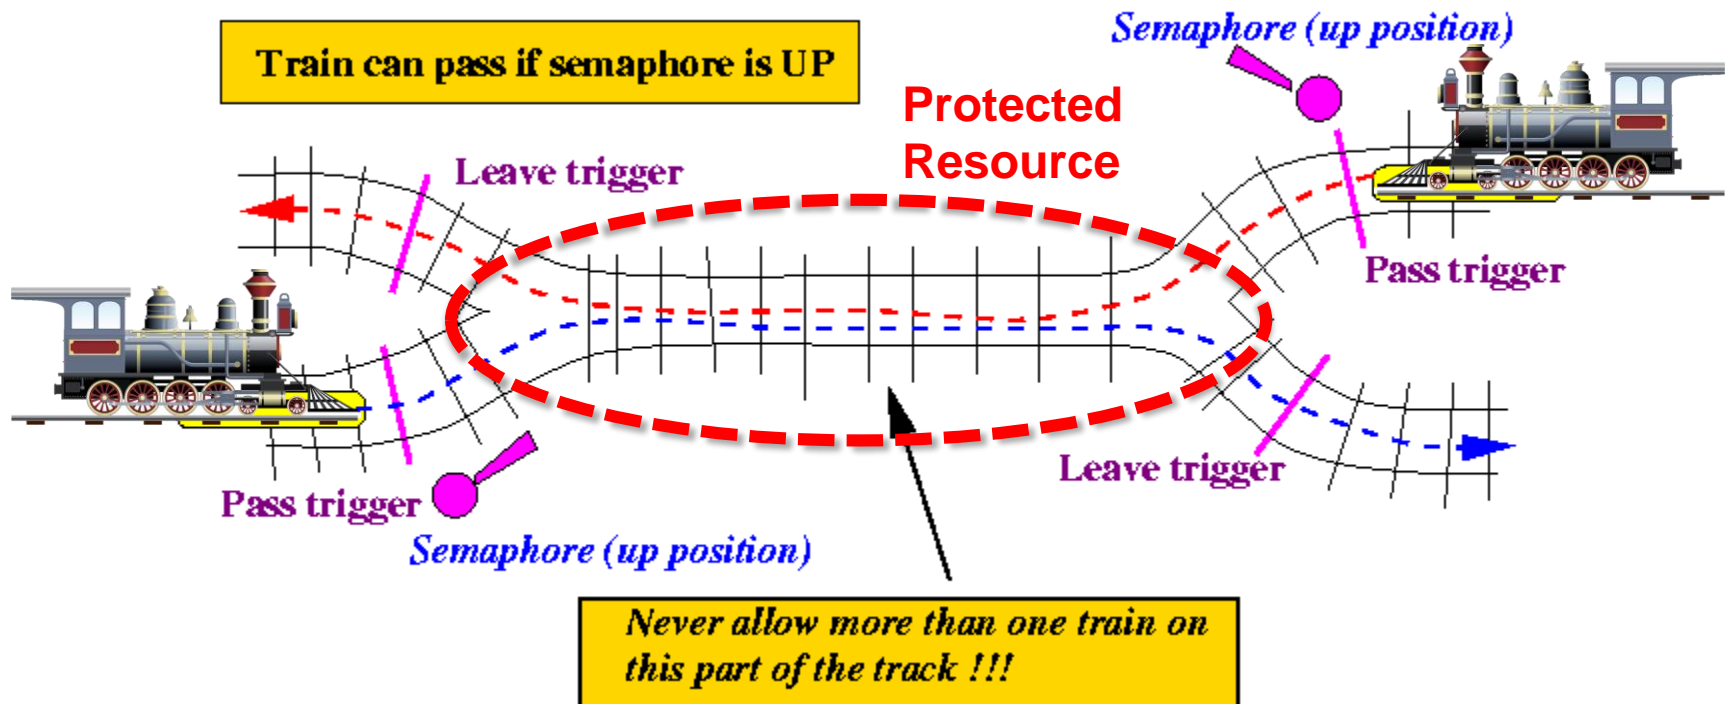


Semaphore Usage

- **Binary semaphore**
- Integer **value** initialized and capped to 1
 - 0 means **lock is held**
 - Same as a lock, but no busy waiting
 - `wait()` is `acquire()`
 - `signal()` is `release()`
- **Counting semaphore**
- Integer **value** initialized and capped to the # of protected resources
 - **Positive** value means still resource available
 - **Negative or zero** value, how many waiting for resources, no resources available
- **Synchronization semaphore**
- Integer **value** initialized to 0
 - **0** means no events pending
 - **Positive** value means there are events pending

Binary Semaphores in Real-life

A Semaphore with **one** resource



Counting Semaphores in Real-life

MAX PEOPLE IN MUSEUM = N

Threads
Processes

Protected
Resource

MUSEUM
ENTRANCE

MUSEUM
EXIT



Synchronization Semaphore

- P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “**synch**” initialized to 0

P1 :

S_1 ;

signal (synch) ;

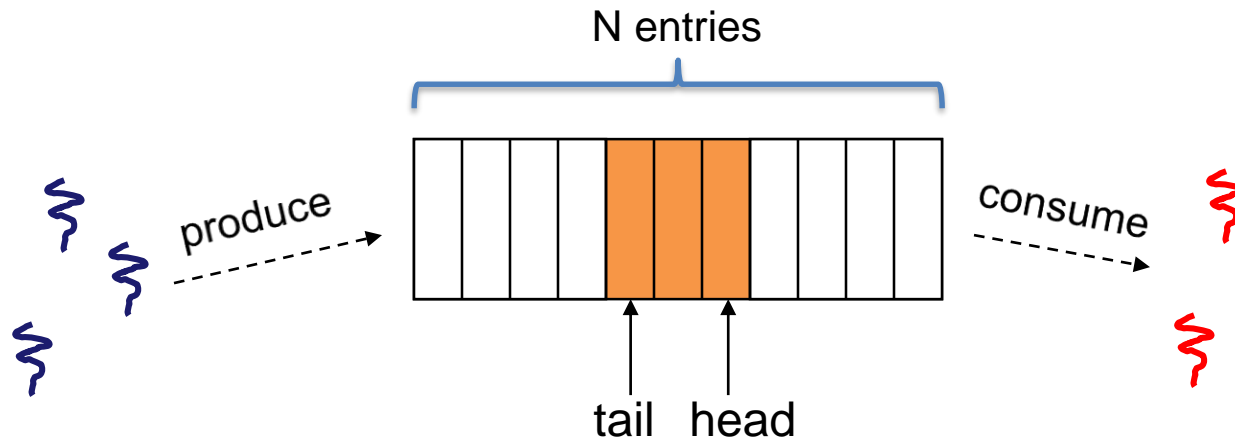
P2 :

wait (synch) ;

S_2 ;

Producer-consumer Problem

- *AKA* **Bounded-buffer**
 - there is a circular buffer in memory with N entries (slots)
 - producer threads insert entries into it (one at a time)
 - consumer threads remove entries from it (one at a time)
- Threads/process are concurrent/parallel
 - Must use synchronization constructs to control access to shared variables describing buffer state



Producer-consumer Problem

Counting Semaphores (1)

Shared variables
`const int N=100;`
`int count=0;`

Previous code



Shared variables
`const int N=100;`
`semaphore empty=N, full=0;`

New code

- In producer-consumer problem ***count*** controls two conditions
 - Buffer empty
 - `count == 1` or `count == 0`
 - Buffer full
 - `count == N` or `count == N-1`
- A semaphore controls a single condition on ***value***
 - Two semaphores are needed

Producer-consumer Problem

Counting Semaphores (2)

Shared variables

```
const int N=100;  
int count=0;
```

Producer

```
while (1){  
    produce an item A;  
    if(count==N) sleep();  
    insert item;  
    count++;  
    if(count==1) wakeup(consumer);  
}
```

Consumer

```
while (1){  
    if(count==0) sleep();  
    remove item;  
    count--;  
    if(count==N-1) wakeup(producer);  
    consume an item;  
}
```

Full?

Empty?

Previous code

Shared variables

```
const int N=100;  
semaphore full=N, empty=0;
```

Producer

```
while (1){  
    produce an item A;  
    wait(full);  
    insert item;  
    signal(empty);  
}
```

Consumer

```
while (1){  
    wait(empty);  
    remove item;  
    signal(full);  
    consume an item;  
}
```

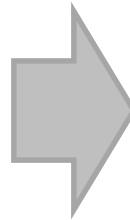
New code

Producer-consumer Problem

Counting and Binary Semaphores (1)

Shared variables
`const int N=100;`
`int count=0;`

Previous code



Shared variables
`const int N=100;`
`semaphore empty=N, full=0;`
`semaphore mux=1;`

New code

- Multiple threads/process trying to add/remove different items on the buffer
 - Operations on the buffer are critical operation
- Mutex to protect buffer operations
 - Avoid races

Producer-consumer Problem

Counting and Binary Semaphores (2)

Shared variables

```
const int N=100;  
int count=0;
```

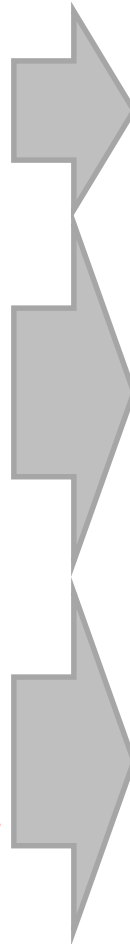
Producer

```
while (1){  
    produce an item A;  
    if(count==N) sleep();  
    insert item;  
    count++;  
    if(count==1) wakeup(consumer);  
}
```

Consumer

```
while (1){  
    if(count==0) sleep();  
    remove item;  
    count--;  
    if(count==N-1) wakeup(producer);  
    consume an item;  
}
```

sleep()&wakeup() code



Shared variables

```
const int N=100;  
semaphore full=N, empty=0;  
semaphore mux=1;
```

Producer

```
while (1){  
    produce an item A;  
    wait(full);  
    wait(mux); // acquire(mux)  
    insert item;  
    signal(mux); // release(mux)  
    signal(empty);  
}
```

Consumer

```
while (1){  
    wait(empty);  
    wait(mux); // acquire(mux)  
    remove item;  
    signal(mux); // release(mux)  
    signal(full);  
    consume an item;  
}
```

New code

Example: Readers/Writers

- Description:
 - A single object is shared among several threads/processes
 - Sometimes a thread just reads the object
 - Sometimes a thread updates (writes) the object
 - **We can allow multiple readers at a time**
 - **Do not change state – no race condition**
 - **We can only allow one writer at a time**
 - Change state- race condition



Readers/Writers Using Semaphores

```
var mutex: semaphore = 1    ; controls access to readcount
    wrt: semaphore = 1    ; control entry for a writer or first reader
    readcount: integer = 0    ; number of active readers
```

```
writer:
    P(wrt)                ; any writers or readers?
    <perform write operation>
    V(wrt)                ; allow others
```

```
reader:
    P(mutex)                ; ensure exclusion
    readcount++            ; one more reader
    if readcount == 1 then P(wrt)    ; if we're the first, synch with writers
    V(mutex)
    <perform read operation>
    P(mutex)                ; ensure exclusion
    readcount--            ; one fewer reader
    if readcount == 0 then V(wrt)    ; no more readers, allow a writer
    V(mutex)
```

Readers/Writers Notes

- Notes
 - the first reader blocks on $P(wrt)$ if there is a writer
 - any other readers will then block on $P(mutex)$
 - if a waiting writer exists, the last reader to exit signals the waiting writer
 - A new reader cannot get in while a writer is waiting
 - When writer exits, if there is both a reader and writer waiting, which one goes next?

Problems with Semaphores and Locks

- Solve any of the **traditional synchronization problems**
- But it is easy to make mistakes
 - Like shared global variables
 - Can be accessed from anywhere (bad software engineering)
 - No connection between the synchronization variable, and the data being controlled
 - No control over their use, no guarantee of proper usage
 - Semaphores: will there ever be a V()?
 - Locks: did you lock when necessary? Unlock at the right time? At all?
- Prone to bugs
 - We can reduce the chance of bugs by “stylizing” the use of synchronization
 - Language help is useful for this



Problem Example

- The sequence of wait and signal operations
 - Producer – wait(full), wait(mux), signal(mux), signal(empty)
 - Consumer – wait(empty), wait(mux), signal(mux), signal(full)
- **Is order of *wait/signal* important?**

Starts with full=N, empty=0, mux=1

Producer

```
while (1){  
  4 produce an item A;  
  5 wait(full);  
  6 wait(mux);  
  7 Blocked!  
  insert item;  
  signal(mux);  
  signal(empty);  
}
```

Consumer

```
while (1){  
  1 wait(mux);  
  2 wait(empty);  
  3 Blocked!  
  remove item;  
  signal(mux);  
  signal(full);  
  consume an item;  
}
```

Both may block forever!

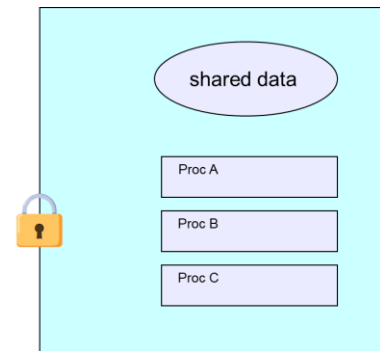
Semaphore Summary

- Solve lock problems
 - No busy-waiting
 - Waiting is controlled by the scheduler
 - Improved scheduling
 - No wasted CPU time
 - Respect task priorities
 - Control multiple resources
- Different usages
 - Binary
 - Counting
 - Synchronization
- Easy to introduce bugs

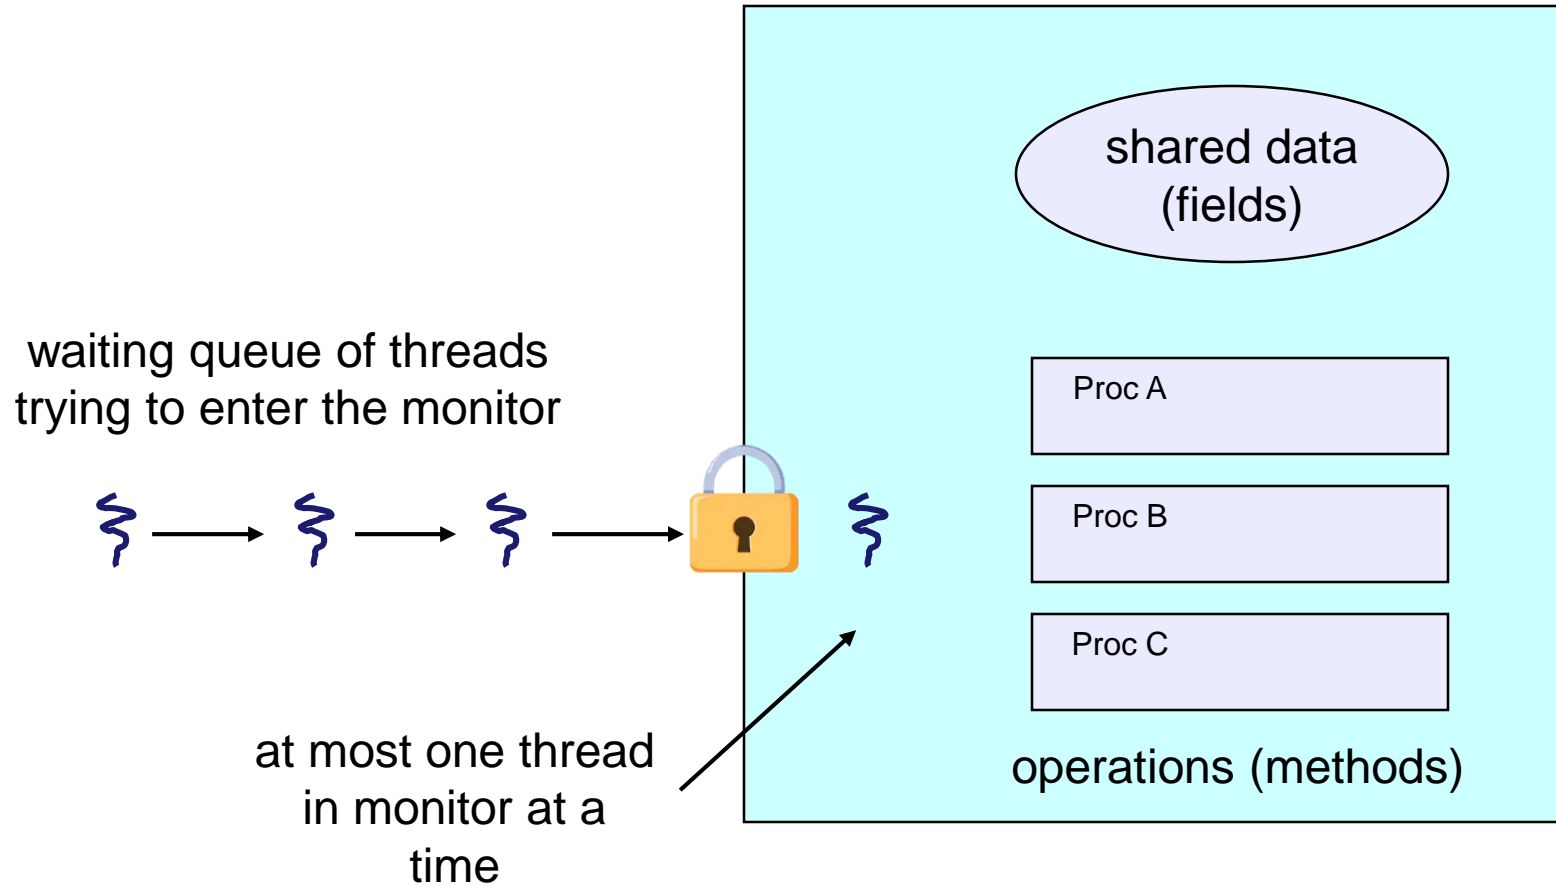
Monitors

Monitors

- Address the key usability issues with semaphores
- A **programming language construct** to support **controlled shared data access**
 - Synchronization code is added by the compiler/language VM
- An **abstract data type/class** in which every method **automatically**
 - acquires a lock on entry
 - releases the lock on exit
- Includes
 - **shared data** structures (object fields)
 - **procedures** that operate on the shared data (object methods)
 - **synchronization** between concurrent execution flows that invoke those procedures
- Data can only be accessed **from within**
 - Protects the data from unstructured access
 - Prevents ambiguity about synchronization



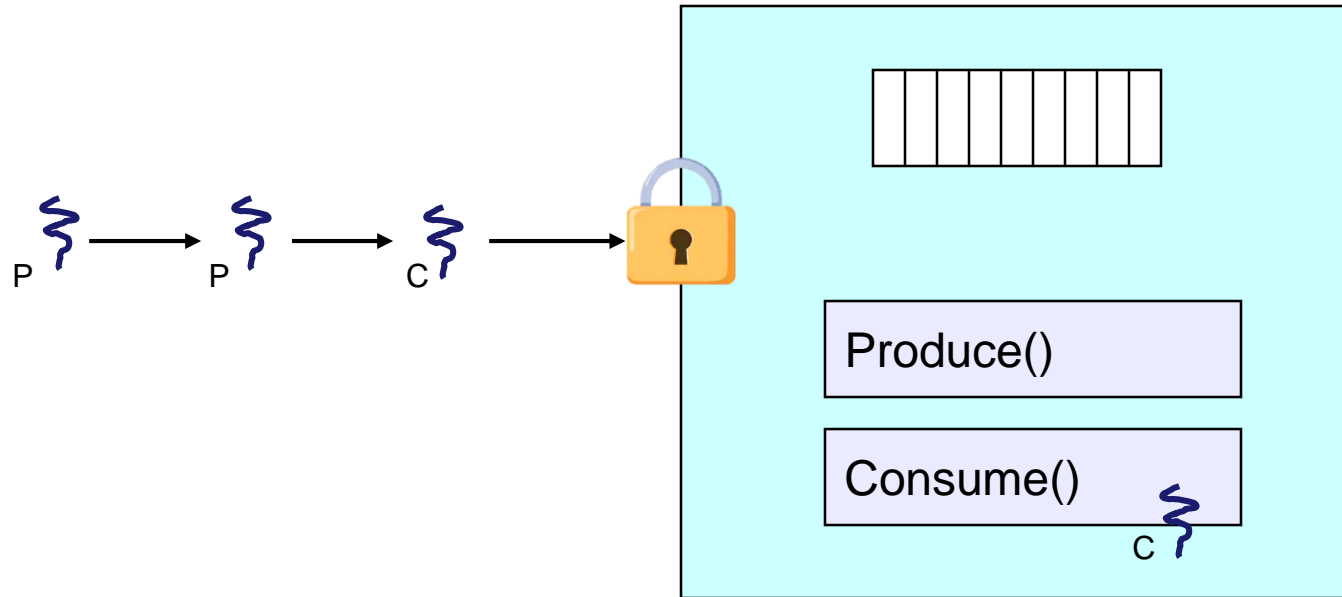
A monitor



Monitor facilities

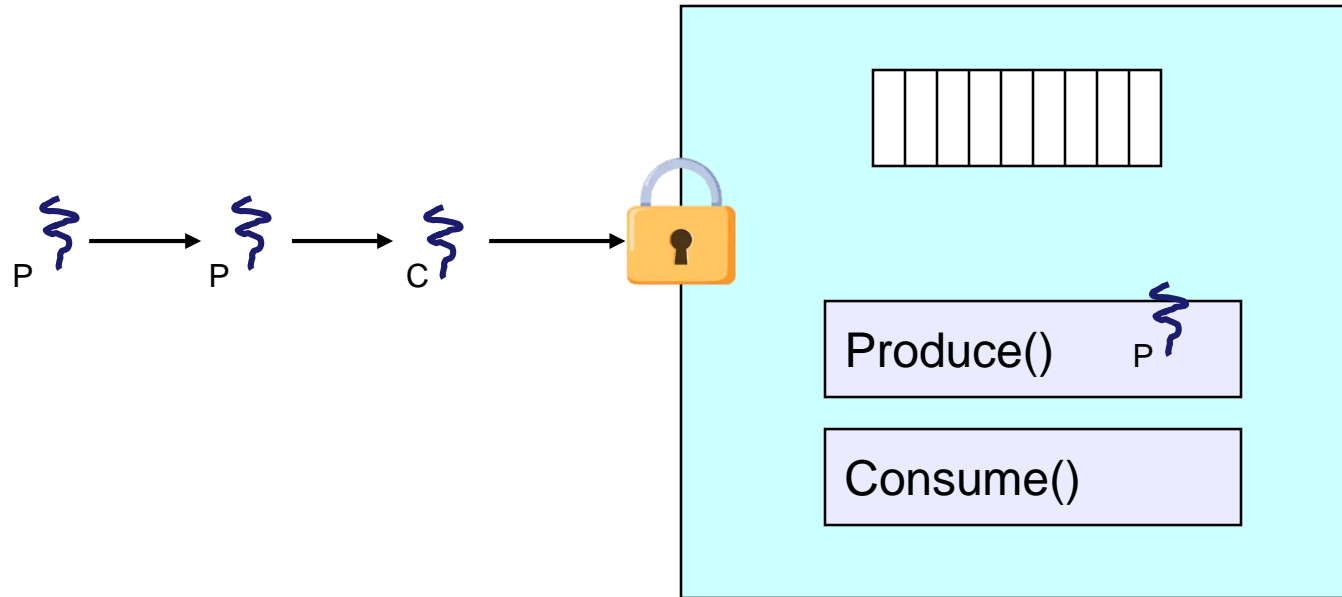
- **Automatic** mutual exclusion
 - Only **one thread** can be executing inside at any time
 - Synchronization is implicitly associated with the monitor
 - it “comes for free”
 - If a second thread tries to execute a monitor procedure it **blocks** until the first has left the monitor
 - More restrictive than semaphores
 - Easier to use (most of the time)
- However, there’s a problem...

Problem: Producer-consumer Scenario



- Buffer is **empty**
- Now what?

Problem: Producer-consumer Scenario



- Buffer is **full**
- Now what?

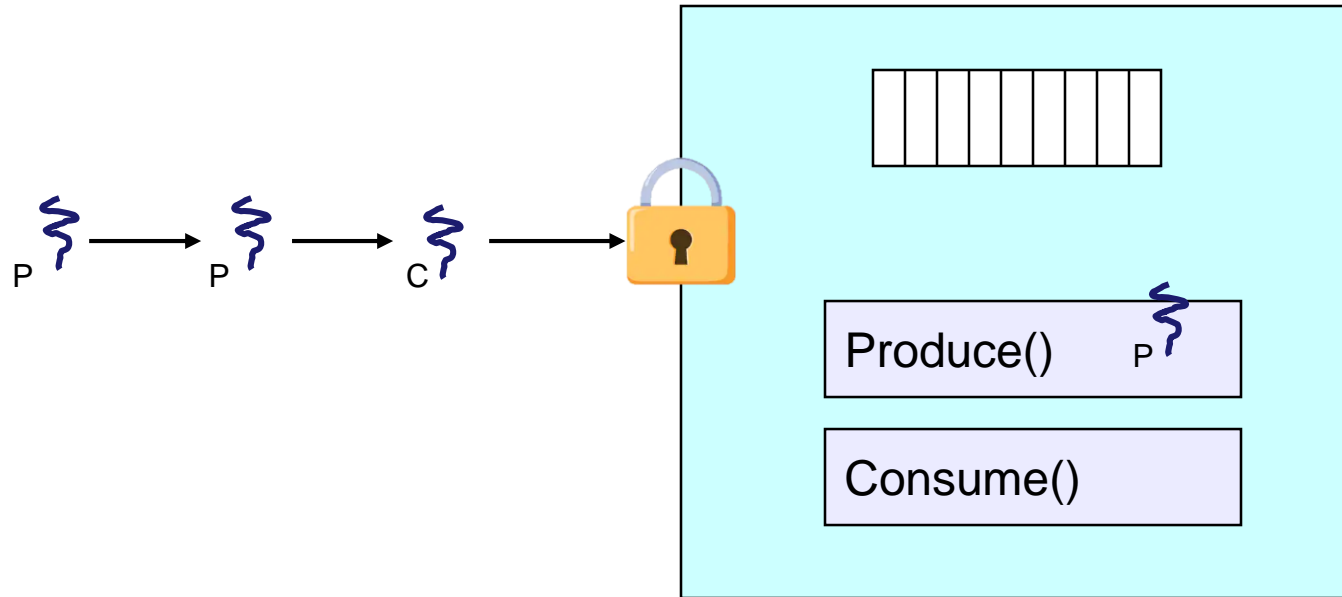
Solution?

- **Condition variables**
- Operations on condition variables
 - **wait(c)**
 - Release monitor lock, so somebody else can get in
 - Wait for somebody else to signal condition
 - Condition variables have associated wait queues
 - **signal(c)**
 - Wake up at most one waiting thread
 - “Hoare” monitor: wakeup immediately, signaller steps outside
 - If no waiting threads, signal is lost
 - this is different than semaphores: no history!
 - **broadcast(c)**
 - Wake up all waiting threads

Producer-consumer using (Hoare) monitors

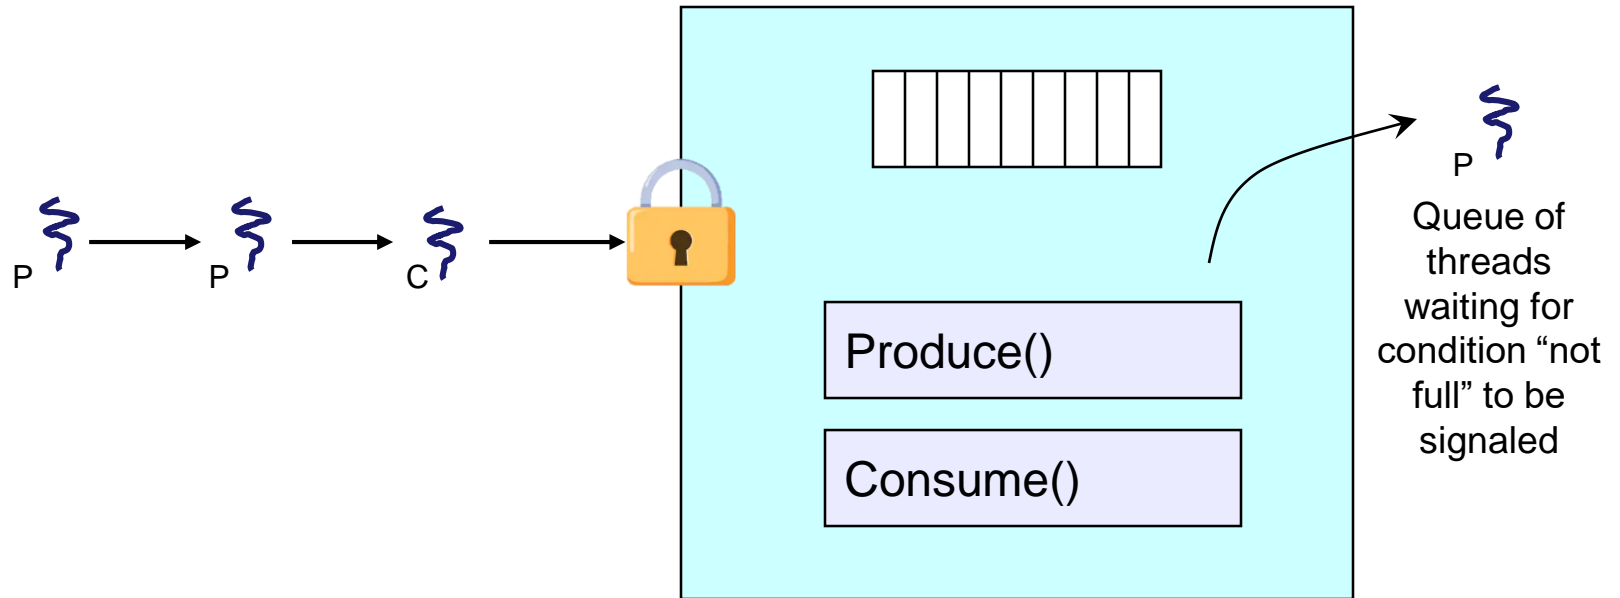
```
Monitor producer_consumer {  
    buffer resources[N];  
    condition not_full, not_empty;  
  
    produce(resource x) {  
        if (array "resources" is full, determined maybe by a count)  
            wait(not_full);  
        insert "x" in array "resources"  
        signal(not_empty);  
    }  
  
    consume(resource *x) {  
        if (array "resources" is empty, determined maybe by a count)  
            wait(not_empty);  
        *x = get resource from array "resources"  
        signal(not_full);  
    }  
}
```

Problem: Producer-consumer Scenario



- Buffer is **full**
- Now what?

Producer-consumer Scenario with Conditional Variable



- Buffer is **full**
- Now what?

Runtime Functions for (Hoare) Monitors

- EnterMonitor(m)
 - guarantee mutual exclusion
- ExitMonitor(m)
 - hit the road, letting someone else run
- Wait(c)
 - step out until condition satisfied
- Signal(c)
 - if someone's waiting, step out and let him run
- EnterMonitor and ExitMonitor are inserted automatically by the compiler
- This guarantees mutual exclusion for code inside of the monitor

Producer-consumer Using (Hoare) Monitors

```
Monitor producer_consumer {  
    buffer resources[N];  
    condition not_full, not_empty;  
  
    produce(resource x) {  
        ..... EnterMonitor(m)  
        if (array "resources" is full, determined maybe by a count)  
            wait(not_full);  
        insert "x" in array "resources"  
        signal(not_empty);  
        ..... ExitMonitor(m)  
    }  
  
    consume(resource *x) {  
        ..... EnterMonitor(m)  
        if (array "resources" is empty, determined maybe by a count)  
            wait(not_empty);  
        *x = get resource from array "resources"  
        signal(not_full);  
        ..... ExitMonitor(m)  
    }  
}
```

Monitor Summary

- Language supported
- Compiler understands them
 - Compiler inserts calls to
 - monitor entry
 - monitor exit
 - Programmer inserts calls to
 - signal
 - wait
 - Language/object encapsulation ensures correctness
 - With conditions, you *still* need to think about synchronization
- Runtime system implements these routines
 - moves threads on and off queues
 - *ensures mutual exclusion!*